
Two LLMs Debate, Both Are Certain They’ve Won

Anonymous Author(s)

Affiliation

Address

email

Abstract

Can LLMs accurately adjust their confidence when facing opposition? Building on previous studies measuring calibration on static fact-based question-answering tasks, we evaluate LLMs in a **dynamic, adversarial debate setting**, uniquely combining two realistic factors: (a) a **multi-turn format** requiring models to update beliefs as new and potentially conflicting information emerges, and (b) a **zero-sum structure** to control for task-related uncertainty, since mutual high-confidence claims (with probabilities summing over 100%) imply systematic overconfidence. We organized 60 three-round policy debates (opening statements, rebuttals, and finals) among ten state-of-the-art LLMs. In each round, models privately rated their confidence (0–100) in winning and explained their reasoning in a hidden scratchpad. We observed five concerning patterns: (1) **Systematic overconfidence**: models began debates overly sure of victory (average initial confidence 72.9% vs. a rational 50% baseline given equal opponents). (2) **Confidence escalation**: rather than reducing confidence toward 50% as debates progressed, both debaters *increased* their win probabilities, averaging 83% by the final round and violating Bayesian updating norms. (3) **Mutual overestimation**: in 61.7% of debates, both sides simultaneously claimed $\geq 75\%$ probability of victory, logically impossible in a zero-sum debate. (4) **Persistent bias in self-debates**: models debating an identical copy increased their confidence from 64.1% to 75.2%; even when explicitly informed their chance of winning was exactly 50%, confidence still rose (from 50.0% to 57.1%). (5) **Misaligned private reasoning**: models’ private scratchpad thoughts often differed from their public confidence ratings, raising concerns about the faithfulness of their reasoning process in strategic settings. These results suggest LLMs possess a fundamental metacognitive flaw, especially evident in realistic multi-turn interactions involving belief updates and controlled uncertainty. This flaw threatens LLM reliability in high-stakes, multi-agent scenarios requiring accurate self-assessment.

1 Introduction

Large language models are increasingly being used in high stakes domains like legal analysis, writing and as agents in deep research Handa et al. [2025] Zheng et al. [2025] which require critical thinking, analysis of competing positions, and iterative reasoning under uncertainty. A foundational skill underlying all of these is calibration—the ability to align one’s confidence with the correctness of one’s beliefs or outputs. In these domains, poorly calibrated confidence can lead to serious errors - an overconfident legal analysis might miss crucial counterarguments, while an uncalibrated research agent might pursue dead ends without recognizing their diminishing prospects. However, language models are often unable to express their confidence in a meaningful or reliable way. While recent work has explored LLM calibration in static, single-turn settings like question answering [Tian et al., 2023, Xiong et al., 2024, Kadavath et al., 2022], real-world reasoning—especially in critical domains like research and analysis—is rarely static or isolated.

Models must respond to opposition, revise their beliefs over time, and recognize when their position is weakening. Their difficulty with introspection and confidence revision in dynamic settings fundamentally limits their usefulness in deliberative settings and poses substantial risks in domains requiring careful judgment under uncertainty. Debate provides a natural framework to stress-test these metacognitive abilities because it requires participants to respond to direct challenges, adapt to new information, and continually reassess the relative strength of competing positions—particularly when their arguments are directly contradicted or new evidence emerges. In adversarial settings, where one side must ultimately prevail, a rational agent should recognize when its position has been weakened and adjust its confidence accordingly. This is especially true when debaters have equal capabilities, as neither should maintain an unreasonable expectation of advantage.

In this work, we study how well language models revise their confidence when engaged in adversarial debate—a setting that naturally stresses the metacognitive abilities crucial for high-stakes applications. We simulate 60 three-round debates between ten state-of-the-art LLMs across six global policy motions. After each round—opening, rebuttal, and final—models provide private, incentivized confidence bets (0-100) estimating their probability of winning, along with natural language explanations in a private scratchpad. The debate setup ensures both sides have equal access to information and equal opportunity to present their case.

Our results reveal a fundamental metacognitive deficit. Key findings include: (1) systematic overconfidence (average opening stated confidence of 72.92% vs. an expected 50% win rate); (2) a pattern of "confidence escalation," where average confidence increased from opening (72.9%) to closing rounds (83.3%), contrary to Bayesian principles, even for losing models; (4) persistent overconfidence even when models debated identical counterparts even though all models know they face opponents of equal capability, with no inherent advantage. In 71.7% of debates, both debaters report high confidence ($\geq 75\%$)—a logically incoherent outcome and (5) misalignment between models' internal assessment and expressed confidence, raising concerns about the faithfulness of chain-of-thought reasoning.

The challenge of LLM calibration becomes particularly acute in dynamic, interactive settings, raising serious concerns about deploying them in roles requiring accurate self-assessment and real-time adaptation to new evidence. We investigate a core aspect of this problem, identifying a pattern we term confidence escalation: an anti-Bayesian drift where LLMs not only systematically overestimate their correctness but often become more certain after facing counter-arguments. This metacognitive blind spot, persistent even when incentives are aligned with accurate self-assessment, threatens reliability in adversarial, multi-agent, and safety-critical applications. For instance, an overconfident LLM might provide flawed legal advice without appropriate caveats, mismanage critical infrastructure in an automated system, or escalate unproductive arguments in collaborative research settings. Until models can reliably revise their confidence in response to opposition, their epistemic judgments in adversarial contexts cannot be trusted—a critical limitation for systems meant to engage in research, analysis, or high-stakes decision making

To probe these critical metacognitive issues, this paper makes several contributions. First, and central to our investigation, we introduce a novel and highly accessible debate-based methodology for studying dynamic confidence calibration in LLMs. A key innovation of our framework is its **self-contained design: it evaluates the coherence and rationality of confidence revisions directly from model interactions, obviating the need for external human judges to assess argument quality or predefined 'ground truth' debate outcomes.** This streamlined approach makes the study of LLM metacognition more scalable and broadly applicable. Second, employing this methodology, we systematically quantify significant overconfidence and the aforementioned confidence escalation phenomenon across various LLMs and debate conditions. Our analysis includes novel findings on model behavior in identical-model debates and the impact of public versus private confidence reporting. Collectively, these contributions highlight fundamental limitations in current LLM self-assessment capabilities, offering crucial insights for AI safety and the responsible development of more epistemically sound AI systems

2 Related Work

Confidence Calibration in LLMs. Recent work has explored methods for eliciting calibrated confidence from large language models (LLMs). While pretrained models have shown relatively

well-aligned token-level probabilities [Kadavath et al., 2022], calibration tends to degrade after reinforcement learning from human feedback (RLHF). To address this, Tian et al. [2023] propose directly eliciting *verbalized* confidence scores from RLHF models, showing that they outperform token probabilities on factual QA tasks. Xiong et al. [2024] benchmark black-box prompting strategies for confidence estimation across multiple domains, finding moderate gains but persistent overconfidence. However, these studies are limited to static, single-turn tasks. In contrast, we evaluate confidence in a multi-turn, adversarial setting where models must update beliefs in response to opposing arguments.

LLM Metacognition and Self-Evaluation. A related line of work examines whether LLMs can reflect on and evaluate their own reasoning. Song et al. [2025] show that models often fail to express knowledge they implicitly encode, revealing a gap between internal representation and surface-level introspection. Other studies investigate post-hoc critique and self-correction Li et al. [2024], but typically focus on revising factual answers, not tracking relative argumentative success. Our work tests whether models can *dynamically monitor* their epistemic standing in a debate—arguably a more socially and cognitively demanding task.

Debate as Evaluation and Oversight. Debate has been proposed as a mechanism for AI alignment, where two agents argue and a human judge evaluates which side is more truthful or helpful [Irving et al., 2018]. More recently, Brown-Cohen et al. [2023] propose “doubly-efficient debate,” showing that honest agents can win even when outmatched in computation, if the debate structure is well-designed. While prior work focuses on using debate to elicit truthful outputs or train models, we reverse the lens: we use debate as a testbed for evaluating *epistemic self-monitoring*. Our results suggest that current LLMs, even when incentivized and prompted to reflect, struggle to track whether they are being outargued.

Persuasion, Belief Drift, and Argumentation. Other studies examine how LLMs respond to external persuasion. Xu et al. [2023] show that models can abandon correct beliefs when exposed to carefully crafted persuasive dialogue. Zhou et al. [2023a] and Rivera et al. [2023] find that language assertiveness influences perceived certainty and factual accuracy. While these works focus on belief change due to stylistic pressure, we examine whether models *recognize when their own position is deteriorating*, and how that impacts their confidence. We find that models often fail to revise their beliefs, even when presented with strong, explicit opposition.

Human Overconfidence Baselines We observe that LLM overconfidence patterns parallel established human cognitive biases. We will discuss and compare existing research on both human and LLM overconfidence in detail in the Discussion section (§??).

Summary. Our work sits at the intersection of calibration, metacognition, adversarial reasoning, and debate-based evaluation. We introduce a new diagnostic setting—structured multi-turn debate with private, incentivized confidence betting—and show that LLMs frequently overestimate their standing, fail to adjust, and exhibit “confidence escalation” despite losing. These findings surface a deeper metacognitive failure that challenges assumptions about LLM trustworthiness in high-stakes, multi-agent contexts.

3 Methodology

Our study investigates the dynamic metacognitive abilities of Large Language Models (LLMs)—specifically their confidence calibration and revision—through a novel experimental paradigm based on competitive policy debate. The primary data for assessing metacognition was gathered via **round-by-round private confidence elicitation**, where models provided a numerical confidence bet (0-100) on their victory and explained their reasoning in a **private scratchpad** after each speech. This allowed us to directly observe their internal self-assessments and their evolution during debate.

To probe these metacognitive behaviors under various conditions, we conducted experiments in **four distinct configurations**:

- 143 1. **Cross-Model Debates:** We conducted 60 debates between different pairs of ten state-of-the-
 144 art LLMs across six policy topics (details on models, topics, and pairings in Appendices A, E
 145 B). These debates provided a general competitive setting to observe how confidence behaves
 146 in heterogeneous matchups. For these debates, where the true outcome was unknown a
 147 priori, an AI jury was employed to provide an external adjudication of win/loss records,
 148 enabling analysis of external calibration (details on jury in Appendix D.4).
- 149 2. **Standard Self-Debates (Jury-Independent Test):** In this configuration, designed for jury-
 150 independent analysis, each of our ten LLMs debated an identical copy of itself across the six
 151 topics. The prompt explicitly stated they were facing an equally capable opponent (details
 152 in Appendix F). This isolated the assessment of internal confidence under known perfect
 153 symmetry and a theoretically 50% win probability, without external judgment.
- 154 3. **Informed Self-Debates (Anchoring Test):** Building on the standard self-debate, models
 155 were additionally and explicitly informed that they had exactly a fifty percent chance
 156 of winning (details in Appendix G). This experiment investigated the influence of direct
 157 probabilistic anchoring on confidence calibration in a jury-independent setting.
- 158 4. **Public Self-Debates (Strategic Signaling Test):** In this configuration, models faced an
 159 identical opponent, were told of the 50% win probability, and crucially, their confidence
 160 bets were made **public** to their opponent (details in Appendix H). This explored the impact
 161 of strategic considerations on reported confidence, providing insight into the faithfulness of
 162 expressed beliefs in a public scenario, also in a jury-independent context for the internal
 163 belief vs. public report comparison.

164 Each configuration involved debates across the six policy topics, with models rotating roles and
 165 opponents as appropriate for the design. The following sections detail the common elements of the
 166 debate setup and the specific analysis conducted for each experimental configuration.

167 3.1 Debate Simulation Environment

168 **Debater Pool:** We utilized ten LLMs, selected to represent diverse architectures and leading providers
 169 (and depicted visually in Figure ??ix A for the full list). In each debate, two models were randomly
 170 assigned to the Proposition and Opposition sides according to a balanced pairing schedule designed
 171 to ensure each model debated a variety of opponents across different topics (see Appendix B for
 172 details).

173 **Debate Topics:** Debates were conducted on six complex global policy motions adapted from the
 174 World Schools Debating Championships corpus. To ensure fair ground and clear win conditions,
 175 motions were modified to include explicit burdens of proof for both sides (see Appendix E for the
 176 full list).

177 3.2 Structured Debate Framework

178 To focus LLMs on substantive reasoning and minimize stylistic variance, we implemented a highly
 179 structured three-round debate format (Opening, Rebuttal, Final).

180 **Concurrent Opening Round:** A key feature of our design was a non-standard opening round where
 181 both Proposition and Opposition models generated their opening speeches simultaneously, based only
 182 on the motion and their assigned side, *before* seeing the opponent’s case. This crucial step allowed
 183 us to capture each LLM’s baseline confidence assessment prior to any interaction or exposure to
 184 opposing arguments.

185 **Subsequent Rounds:** Following the opening, speeches were exchanged, and the debate proceeded
 186 through a Rebuttal and Final round. When generating its speech in these subsequent rounds, each
 187 model had access to the full debate history from all preceding rounds (e.g., for the Rebuttal, both
 188 Opening speeches were available; for the Final, both Opening and both Rebuttal speeches were
 189 available). However, to maintain the symmetrical information state established in the simultaneous
 190 opening and avoid giving either side an immediate preview advantage within a round, neither the
 191 Proposition nor the Opposition model saw the opponent’s speech for that specific round (e.g., the
 192 opponent’s Rebuttal) before generating their own. Both models formulated their arguments based
 193 on the cumulative case presented in the history up to the start of that round, rather than as direct,

194 real-time responses to the opponent’s points in that turn. This design allowed us to evaluate how
195 models integrated and responded to the opponent’s case as it built over time, while ensuring fairness.

196 3.3 Core Prompt Structures & Constraints

197 Highly structured prompts were used for *each* speech type to ensure consistency and enforce specific
198 argumentative tasks, thereby isolating reasoning and self-assessment capabilities. The core structure
199 and key required components for the Opening, Rebuttal, and Final speech prompts are illustrated in
200 Figure 1.

201 Highly structured prompts were used for *each* speech type to ensure consistency and enforce specific
202 argumentative tasks, thereby isolating reasoning and self-assessment capabilities.

203 **Embedded Judging Guidance:** Crucially, all debater prompts included explicit **Judging Guidance** ,
204 instructing debaters on the importance of direct clash, evidence quality hierarchy, logical validity,
205 response obligations, and impact analysis, while explicitly stating that rhetoric and presentation style
206 would be ignored.

207 Full verbatim prompt text for debaters is provided in Appendix C.

208 3.4 Dynamic Confidence Elicitation

209 After generating the content for *each* of their three speeches (including the concurrent opening),
210 models were required to provide a private “confidence bet”.

211 **Mechanism:** This involved outputting a numerical value from 0 to 100, representing their perceived
212 probability of winning the debate, using a specific XML tag (<bet_amount>). Models were also
213 prompted to provide private textual justification for their bet amount within separate XML tags
214 (<bet_logic_private>), allowing for qualitative insight into their reasoning.

215 **Purpose:** This round-by-round elicitation allowed us to quantitatively track self-assessed performance
216 dynamically throughout the debate, enabling analysis of confidence levels, calibration, and revision
217 (or lack thereof) in response to the evolving argumentative context.

218 3.5 Data Collection

219 The final dataset comprises the full transcripts of 240 debates, the round-by-round confidence bets
220 (amount and private thoughts) from both debaters in each debate, and the detailed structured verdicts
221 (winner, confidence, reasoning) from each of the six AI judges for the cross-model debates. This data
222 enables the quantitative analysis of LLM overconfidence, confidence revision and calibration for the
223 cross-model debates presented in our findings.

224 4 Results

225 Our experimental setup, involving 60 simulated policy debates per configuration between ten state-
226 of-the-art LLMs, with round-by-round confidence elicitation yielded several key findings regarding
227 LLM metacognition in adversarial settings.

228 4.1 Pervasive Overconfidence Without Seeing Opponent Argument (Finding 1)

229 A core finding across all four experimental configurations was significant LLM overconfidence,
230 particularly evident in the initial concurrent opening round before models had seen any counterargu-
231 ments. Given the inherent nature of a two-participant debate where one side wins and the other loses,
232 a rational model should assess its baseline probability of winning at 50% anticipating that the other
233 debater too would make good arguments; however, observed initial confidence levels consistently
234 and substantially exceeded this expectation.

235 As shown in Table 1, the overall average initial confidence reported by models in the Cross-model,
236 Standard Self, and Public Bets configurations was consistently and significantly above the 50%
237 baseline. Specifically, the mean initial confidence was 72.92% (± 7.93 SD, $n=120$) for Cross-
238 model debates, 64.08% (± 15.32 SD, $n=120$) for Standard Self debates (private bets without 50%

```

===== OPENING SPEECH PROMPT =====

ARGUMENT 1
Core Claim: (State your first main claim in one clear sentence)
Support Type: (Choose either EVIDENCE or PRINCIPLE)
Support Details:
  For Evidence:
    - Provide specific examples with dates/numbers
    - Include real world cases and outcomes
    - Show clear relevance to the topic
  For Principle:
    - Explain the key principle/framework
    - Show why it is valid/important
    - Demonstrate how it applies here
Connection: (Explicit explanation of how this evidence/principle proves claim)

ARGUMENT 2
(Use exact same structure as Argument 1)

ARGUMENT 3 (Optional)
(Use exact same structure as Argument 1)

SYNTHESIS
- Explain how your arguments work together as a unified case
- Show why these arguments prove your side of the motion
- Present clear real-world impact and importance
- Link back to key themes/principles

JUDGING GUIDANCE (excerpt)
Direct Clash - Evidence Quality Hierarchy - Logical Validity -
Response Obligations - Impact Analysis & Weighing
-----

===== REBUTTAL SPEECH PROMPT =====

CLASH POINT 1
Original Claim: (Quote opponent's exact claim)
Challenge Type: Evidence Critique | Principle Critique |
                Counter Evidence | Counter Principle
Challenge:
  (Details depend on chosen type; specify flaws or present counters)
Impact: (Explain why winning this point is crucial)

CLASH POINT 2, 3 (same template)

DEFENSIVE ANALYSIS
  Vulnerabilities - Additional Support - Why We Prevail

WEIGHING
  Key Clash Points - Why We Win - Overall Impact

JUDGING GUIDANCE (same five criteria as above)
-----

===== FINAL SPEECH PROMPT =====

FRAMING
Core Questions: (Identify fundamentals and evaluation lens)

KEY CLASHES (repeat for each major clash)
Quote: (Exact disagreement)
Our Case Strength: (Show superior evidence/principle)
Their Response Gaps: (Unanswered flaws)
Crucial Impact: (Why this clash decides the motion)

VOTING ISSUES
Priority Analysis - Case Proof - Final Weighing

JUDGING GUIDANCE (same five criteria as above)
=====

```

Figure 1: Structured prompts supplied to LLM debaters for the opening, rebuttal, and final speeches. Full, unabridged text appears in the appendix.

Table 1: Mean (\pm Standard Deviation) Initial Confidence (0-100%) Reported by LLMs Across Experimental Configurations. Sample size (n) per model per configuration is indicated in parentheses. The 'Standard Self' condition represents private bets in self-debates without explicit probability instruction, while 'Informed Self' includes explicit instruction about the 50% win probability.

Model	Cross-model	Standard Self	Informed Self (50% informed)	Public Bets (Public Bets)
anthropic/claude-3.5-haiku	71.67 \pm 4.92 (n=12)	71.25 \pm 6.44 (n=12)	54.58 \pm 9.64 (n=12)	73.33 \pm 7.18 (n=12)
anthropic/claude-3.7-sonnet	67.31 \pm 3.88 (n=13)	56.25 \pm 8.56 (n=12)	50.08 \pm 2.15 (n=12)	56.25 \pm 6.08 (n=12)
deepseek/deepseek-chat	74.58 \pm 7.22 (n=12)	54.58 \pm 4.98 (n=12)	49.17 \pm 6.34 (n=12)	56.25 \pm 7.42 (n=12)
deepseek/deepseek-r1-distill-qwen-14b:free	79.09 \pm 10.44 (n=11)	76.67 \pm 13.20 (n=12)	55.75 \pm 4.71 (n=12)	69.58 \pm 16.30 (n=12)
google/gemini-2.0-flash-001	65.42 \pm 8.38 (n=12)	43.25 \pm 27.03 (n=12)	36.25 \pm 26.04 (n=12)	34.58 \pm 25.80 (n=12)
google/gemma-3-27b-it	67.50 \pm 6.22 (n=12)	68.75 \pm 7.42 (n=12)	53.33 \pm 11.15 (n=12)	63.75 \pm 9.80 (n=12)
openai/gpt-4o-mini	75.00 \pm 3.69 (n=12)	67.08 \pm 7.22 (n=12)	57.08 \pm 12.70 (n=12)	72.92 \pm 4.98 (n=12)
openai/o3-mini	77.50 \pm 5.84 (n=12)	70.00 \pm 10.66 (n=12)	50.00 \pm 0.00 (n=12)	72.08 \pm 9.40 (n=12)
qwen/qwen-max	73.33 \pm 8.62 (n=12)	62.08 \pm 12.87 (n=12)	43.33 \pm 22.29 (n=12)	64.58 \pm 10.97 (n=12)
qwen/qwq-32b:free	78.75 \pm 4.33 (n=12)	70.83 \pm 10.62 (n=12)	50.42 \pm 1.44 (n=12)	71.67 \pm 8.62 (n=12)
OVERALL AVERAGE	72.92 \pm 7.93 (n=120)	64.08 \pm 15.32 (n=120)	50.00 \pm 13.61 (n=120)	63.50 \pm 16.38 (n=120)

instruction), and 63.50% (\pm 16.38 SD, n=120) for Public Bets (public bets without 50% instruction). One-sample t-tests confirmed that the mean initial confidence in each of these three conditions was statistically significantly greater than 50% (Cross-model: $t=31.67$, $p<0.001$; Standard Self: $t=10.07$, $p<0.001$; Public Bets: $t=9.03$, $p<0.001$). Wilcoxon signed-rank tests yielded similar conclusions (all $p<0.001$), confirming the robustness of this finding to distributional assumptions. This pervasive overconfidence in the initial assessment, before any interaction with an opponent's case, suggests a fundamental miscalibration bias in LLMs' self-assessment of their standing in a competitive context.

In stark contrast, the overall average initial confidence in the Informed Self configuration was precisely 50.00% (\pm 13.61 SD, n=120). A one-sample t-test confirmed that this mean was not statistically significantly different from 50% ($t=0.00$, $p=1.0$). Furthermore, a paired t-test comparing the per-model means in the Standard Self and Informed Self configurations revealed a statistically significant reduction in initial confidence when models were explicitly informed of the 50% win probability (mean difference = 14.08, $t=7.07$, $p<0.001$). This demonstrates that while the default state is overconfident, models can align their *initial* reported confidence much closer to the rational baseline when explicitly anchored with the correct probability.

Analysis at the individual model level (see Appendix ?? for full results) shows that this overconfidence was widespread, with 30 out of 40 individual model-configuration combinations showing initial confidence significantly greater than 50% (one-sided t-tests, $\alpha = 0.05$). However, we also observed considerable variability in initial confidence (large standard deviations), both across conditions and for specific models like Google Gemini 2.0 Flash (\pm 27.03 SD in Standard Self). Notably, some models, such as OpenAI o3-Mini and Qwen QWQ-32b, reported perfectly calibrated initial confidence (50.00 \pm 0.00 SD) in the Informed Self condition. The non-significant difference in overall mean initial confidence between Standard Self and Public Bets (mean difference = 0.58, $t=0.39$, $p=0.708$) suggests that simply making the initial bet public does not, on average, significantly alter the self-assessed confidence compared to the private default.

4.2 Confidence Escalation among models (Finding 2)

Building upon the pervasive initial overconfidence (Section 4.1), a second critical pattern observed across *all four* experimental configurations was a significant **confidence escalation**. This refers to the consistent tendency for models' self-assessed probability of winning to increase over the course of the debate, from the initial Opening round to the final Closing statements. As illustrated in Table 2, the overall mean confidence across models rose substantially in every configuration. For instance, mean confidence increased from 72.92% to 83.26% in Cross-model debates, from 64.08% to 75.20% in Standard Self-debates, from 63.50% to 74.15% in Public Bets, and notably, even from a calibrated 50.00% to 57.08% in Informed Self-debates. Paired statistical tests confirmed these overall increases from Opening to Closing were highly significant in all configurations (all $p<0.001$). While this pattern of escalation was statistically significant on average across each configuration, the magnitude and statistical significance of escalation varied at the individual model level (see Appendix K for full per-model test results). This widespread and significant upward drift in self-confidence is highly irrational, particularly evident in the self-debate conditions where models know they face an equally capable opponent and the rational win probability is 50% from the outset. Escalating confidence

in this context, especially when starting near the correct 50% as in the Informed Self condition, demonstrates a fundamental failure to dynamically process adversarial feedback and objectively assess relative standing, defaulting instead to an unjustified increase in self-assurance regardless of the opponent’s performance or the debate’s progression.

Table 2: Overall Mean Confidence (0-100%) and Escalation Across Debate Rounds by Experimental Configuration. Values show Mean \pm Standard Deviation (N). Δ indicates mean change from the earlier to the later round, with paired t-test p-values shown (* $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$).

Experiment Type	Opening Bet	Rebuttal Bet	Closing Bet	Open→Rebuttal	Rebuttal→Closing	Open→Closing
Cross-model	72.92 \pm 7.89 (N=120)	77.67 \pm 9.75 (N=120)	83.26 \pm 10.06 (N=120)	$\Delta=4.75$, $p < 0.001$ ***	$\Delta=5.59$, $p < 0.001$ ***	$\Delta=10.34$, $p < 0.001$ ***
Informed Self	50.00 \pm 13.55 (N=120)	55.77 \pm 9.73 (N=120)	57.08 \pm 8.97 (N=120)	$\Delta=5.77$, $p < 0.001$ ***	$\Delta=1.32$, $p=0.0945$	$\Delta=7.08$, $p < 0.001$ ***
Public Bets	63.50 \pm 16.31 (N=120)	69.43 \pm 16.03 (N=120)	74.15 \pm 14.34 (N=120)	$\Delta=5.93$, $p < 0.001$ ***	$\Delta=4.72$, $p < 0.001$ ***	$\Delta=10.65$, $p < 0.001$ ***
Standard Self	64.08 \pm 15.25 (N=120)	69.07 \pm 16.63 (N=120)	75.20 \pm 15.39 (N=120)	$\Delta=4.99$, $p < 0.001$ ***	$\Delta=6.13$, $p < 0.001$ ***	$\Delta=11.12$, $p < 0.001$ ***
GRAND OVERALL	62.62 \pm 15.91 (N=480)	67.98 \pm 15.57 (N=480)	72.42 \pm 15.71 (N=480)	$\Delta=5.36$, $p < 0.001$***	$\Delta=4.44$, $p < 0.001$***	$\Delta=9.80$, $p < 0.001$***

4.3 Logical Impossibility: Simultaneous High Confidence (Finding 3)

Stemming directly from the observed confidence escalation, we found that LLMs frequently ended debates holding mutually exclusive high confidence in their victory, a mathematically impossible outcome in a zero-sum competition. Specifically, we analyzed the distribution of confidence levels for *both* debate participants in the closing round across all experimental configurations. As summarized in Table 3, a substantial percentage of debates concluded with both models reporting confidence levels of 75% or higher.

Table 3: Distribution of Confidence Level Combinations for Both Debaters in the Closing Round, by Experiment Type. Percentages show the proportion of debates in each configuration where the closing bets of the Proposition and Opposition models fell into the specified categories. The ‘Both >75%’ column represents the core logical inconsistency finding.

Experiment Type	Total Debates	Both $\leq 50\%$	Both 51-75%	Both >75%	50%+51-75%	50%+>75%	51-75%+>75%
cross_model	60	0.0%	6.7%	61.7%	0.0%	0.0%	31.7%
self_debate	60	0.0%	26.7%	35.0%	5.0%	0.0%	33.3%
informed_self	60	23.3%	56.7%	0.0%	15.0%	0.0%	5.0%
public_bets	60	1.7%	26.7%	33.3%	3.3%	1.7%	33.3%
overall	240	6.2%	29.2%	32.5%	5.8%	0.4%	25.8%

In Cross-model debates, a striking **61.7%** ($n = 37/60$) concluded with both the Proposition and Opposition models reporting a confidence of 75% or greater (Table 3, ‘Both >75%’ column). This is a direct manifestation of logical inconsistency at the system level, where the combined self-assessed probabilities of winning drastically exceed the theoretical maximum of 100% for two agents in a zero-sum game.

While less frequent than in the standard Cross-model setting, this logical impossibility was still common in other non-informed configurations. In Standard Self-debates, where models faced an identical twin, 35.0% ($n = 21/60$) showed both participants claiming >75% confidence in the final round. Public Bets debates exhibited a similar rate of simultaneous >75% confidence at 33.3% ($n = 20/60$). The overall rate of this specific logical inconsistency across all 240 non-informed self- and cross-model debates was 32.5% ($n = 78/240$).

Crucially, this type of severe logical inconsistency was entirely absent (0.0%, $n = 0/60$) in the Informed Self configuration. This aligns with our finding that explicit anchoring mitigated initial overconfidence and somewhat reduced the magnitude of subsequent escalation, thereby preventing models from reaching the high, mutually exclusive confidence levels seen in other conditions.

Beyond the most severe ‘Both >75%’ inconsistency, a significant proportion of debates across all configurations saw both participants reporting confidence between 51-75% (overall 29.2%). Combined with the >75% cases, this means that in over 60% of debates (32.5% + 29.2% overall), *both* models finished with confidence above 50%, further illustrating a systemic failure to converge towards a state reflecting the actual debate outcome or the zero-sum nature of the task. The remaining categories in Table 3 indicate scenarios where confidence levels were split across categories, including a small percentage where both models reported low confidence ($\leq 50\%$).

This prevalence of debates ending with simultaneously high confidence directly results from models independently escalating their beliefs without adequately integrating or believing the strength of the opponent’s counterarguments. It reveals a profound disconnect between their internal confidence reporting mechanisms and the objective reality of a competitive, zero-sum task.

4.4 Strategic Confidence in Public Settings (Finding 5)

5 Discussion

[NEW CONTENT THROUGHOUT SECTION 5, TBA]

5.1 Metacognitive Limitations and Possible Explanations

Our findings reveal significant limitations in LLMs’ metacognitive abilities, specifically their capacity to accurately assess their argumentative position and revise confidence in adversarial contexts. Several explanations may account for these observed patterns, including both human-like biases and LLM-specific factors:

Human-like biases

- **Baseline debate overconfidence:** Research on human debaters [?] found that college debate participants estimated their odds of winning at approximately 65% on average, suggesting that high baseline confidence is prevalent for humans in debate settings similar to our experimental design with LLMs.
- **Persistent miscalibration:** Human psychology reveals systematic miscalibration patterns that parallel our findings. Like humans, LLMs exhibit limited accuracy improvement over repeated trials [Moore and Healy, 2008], mirroring our results.
- **Evidence weighting bias:** Crucially, seminal work by Griffin and Tversky [Griffin and Tversky, 1992] found that humans overweight the strength of evidence favoring their beliefs while underweighting its credibility or weight, leading to overconfidence when strength is high but weight is low.
- **Numerical attractor state:** The average LLM confidence ($\sim 73\%$) recalls the human $\sim 70\%$ "attractor state" often used for probability terms like "probably/likely" [Hashim, 2024, Mandel, 2019], potentially a learned artifact of alignment processes that steer LLMs towards human-like patterns [West and Potts, 2025].

LLM-specific factors

- **General overconfidence across models:** Research has shown that LLMs demonstrate systematic overconfidence across various tasks [Chhikara, 2025, Xiong et al., 2024], with larger LLMs exhibiting greater overconfidence on difficult tasks while smaller LLMs show more consistent overconfidence across task types [Wen et al., 2024].
- **RLHF amplification effects:** Post-training for human preferences appears to significantly exacerbate overconfidence. Models trained via RLHF are more likely to indicate high certainty even when incorrect [Leng et al., 2025] and disproportionately output 7/10 for ratings [West and Potts, 2025, OpenAI et al., 2024], suggesting alignment processes inadvertently reinforce confidence biases.
- **Failure to appropriately integrate new evidence:** Wilie et al. [2024] introduced the Belief-R benchmark and showed that most models fail to appropriately revise their initial conclusions after receiving additional, contradicting information. Rather than reducing confidence when they should, models tend to stick to their initial stance. Agarwal and Khanna [2025] found that LLMs can be swayed to believe falsehoods with persuasive, verbose reasoning. Even smaller models can craft arguments that override truthful answers with high confidence, suggesting that LLMs may be susceptible to confident but flawed counterarguments.
- **Training data imbalance:** Training datasets predominantly feature successful task completion rather than explicit failures or uncertainty. This imbalance may limit models’ ability to recognize and represent losing positions accurately [Zhou et al., 2023b].

361 These combined factors likely contribute to the confidence escalation phenomenon we observe, where
362 models fail to properly update their beliefs in the face of opposing arguments.

363 5.2 Implications for AI Safety and Deployment

364 [ADD REFERENCE TO 3.6, PUBLIC VS PRIVATE COT AND IMPLICATIONS ON COT
365 FAITHFULNESS]

366 The confidence escalation phenomenon identified in this study has significant implications for AI
367 safety and responsible deployment. In high-stakes domains like legal analysis, medical diagnosis,
368 or research, overconfident systems may fail to recognize when they are wrong or when additional
369 evidence should cause belief revision.

370 The persistence of overconfidence even in controlled experimental conditions suggests this is a
371 fundamental limitation rather than a context-specific artifact. This has particular relevance for
372 multi-agent systems, where models must negotiate, debate, and potentially admit error to achieve
373 optimal outcomes. If models maintain high confidence despite opposition, they may persist in flawed
374 reasoning paths or fail to incorporate crucial counterevidence.

375 5.3 Potential Mitigations and Guardrails

376 Our ablation study testing explicit 50% win probability instructions shows [placeholder for results].
377 This suggests that direct prompting approaches may help mitigate but not eliminate confidence biases.

378 Other potential mitigation strategies include:

- 379 • Developing dedicated calibration training objectives
- 380 • Implementing confidence verification systems through external validation
- 381 • Creating debate frameworks that explicitly penalize overconfidence or reward accurate
382 calibration
- 383 • Designing multi-step reasoning processes that force models to consider opposing viewpoints
384 before finalizing confidence assessments

385 5.4 Future Research Directions

386 Future work should explore several promising directions:

- 387 • Investigating whether human-LLM hybrid teams exhibit better calibration than either humans
388 or LLMs alone
- 389 • Developing specialized training approaches specifically targeting confidence calibration in
390 adversarial contexts
- 391 • Exploring the relationship between model scale, training methods, and confidence calibration
- 392 • Testing whether emergent abilities in frontier models include improved metacognitive
393 assessments
- 394 • Designing debates where confidence is directly connected to resource allocation or other
395 consequential decisions

396 6 Conclusion

397 — YOUR CONCLUSION CONTENT HERE —

398 References

- 399 Mahak Agarwal and Divyam Khanna. When persuasion overrides truth in multi-agent llm debates:
400 Introducing a confidence-weighted persuasion override rate (cw-por), 2025. URL [https://](https://arxiv.org/abs/2504.00374)
401 arxiv.org/abs/2504.00374.

402 Jonah Brown-Cohen, Geoffrey Irving, and Georgios Piliouras. Scalable ai safety via doubly-efficient
403 debate. *arXiv preprint arXiv:2311.14125*, 2023. URL <https://arxiv.org/abs/2311.14125>.

404 Prateek Chhikara. Mind the confidence gap: Overconfidence, calibration, and distractor effects in
405 large language models, 2025. URL <https://arxiv.org/abs/2502.11028>.

406 Dale Griffin and Amos Tversky. The weighing of evidence and the determinants of confidence.
407 *Cognitive Psychology*, 24(3):411–435, 1992. doi: [https://doi.org/10.1016/0010-0285\(92\)90013-R](https://doi.org/10.1016/0010-0285(92)90013-R).

408 Kunal Handa, Alex Tamkin, Miles McCain, Saffron Huang, Esin Durmus, Sarah Heck, Jared Mueller,
409 Jerry Hong, Stuart Ritchie, Tim Belonax, Kevin K. Troy, Dario Amodei, Jared Kaplan, Jack Clark,
410 and Deep Ganguli. Which economic tasks are performed with ai? evidence from millions of claude
411 conversations, 2025. URL <https://arxiv.org/abs/2503.04761>.

412 Muhammad J. Hashim. Verbal probability terms for communicating clinical risk - a systematic review.
413 *Ulster Medical Journal*, 93(1):18–23, Jan 2024. Epub 2024 May 3.

414 Geoffrey Irving, Paul Christiano, and Dario Amodei. Ai safety via debate. *arXiv preprint*
415 *arXiv:1805.00899*, 2018. URL <https://arxiv.org/abs/1805.00899>.

416 Saurav Kadavath, Tom Conerly, Amanda Askell, Tom Henighan, Dawn Drain, Ethan Perez, Nicholas
417 Schiefer, Zac Hatfield-Dodds, Nova DasSarma, Eli Tran-Johnson, et al. Language models (mostly)
418 know what they know. *arXiv preprint arXiv:2207.05221*, 2022. URL <https://arxiv.org/abs/2207.05221>.

420 Jixuan Leng, Chengsong Huang, Banghua Zhu, and Jiaxin Huang. Taming overconfidence in llms:
421 Reward calibration in rlhf, 2025. URL <https://arxiv.org/abs/2410.09724>.

422 Loka Li, Guan-Hong Chen, Yusheng Su, Zhenhao Chen, Yixuan Zhang, Eric P. Xing, and Kun
423 Zhang. Confidence matters: Revisiting intrinsic self-correction capabilities of large language
424 models. *ArXiv*, abs/2402.12563, 2024. URL <https://api.semanticscholar.org/CorpusID:268032763>.

426 David R. Mandel. Systematic monitoring of forecasting skill in strategic intelligence. In David R.
427 Mandel, editor, *Assessment and Communication of Uncertainty in Intelligence to Support Decision*
428 *Making: Final Report of Research Task Group SAS-114*, page 16. NATO Science and Technol-
429 ogy Organization, Brussels, Belgium, March 2019. URL [https://papers.ssrn.com/sol3/](https://papers.ssrn.com/sol3/papers.cfm?abstract_id=3435945)
430 [papers.cfm?abstract_id=3435945](https://papers.ssrn.com/sol3/papers.cfm?abstract_id=3435945). Posted: 15 Aug 2019, Conditionally accepted.

431 Don A. Moore and Paul J. Healy. The trouble with overconfidence. *Psychological Review*, 115(2):
432 502–517, 2008. doi: <https://doi.org/10.1037/0033-295X.115.2.502>.

433 OpenAI, Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni
434 Aleman, Diogo Almeida, Janko Altenschmidt, Sam Altman, Shyamal Anadkat, Red Avila, Igor
435 Babuschkin, Suchir Balaji, Valerie Balcom, Paul Baltescu, Haiming Bao, Mohammad Bavarian,
436 Jeff Belgum, Irwan Bello, Jake Berdine, Gabriel Bernadett-Shapiro, Christopher Berner, Lenny
437 Bogdonoff, Oleg Boiko, Madelaine Boyd, Anna-Luisa Brakman, Greg Brockman, Tim Brooks,
438 Miles Brundage, Kevin Button, Trevor Cai, Rosie Campbell, Andrew Cann, Brittany Carey, Chelsea
439 Carlson, Rory Carmichael, Brooke Chan, Che Chang, Fotis Chantzis, Derek Chen, Sully Chen,
440 Ruby Chen, Jason Chen, Mark Chen, Ben Chess, Chester Cho, Casey Chu, Hyung Won Chung,
441 Dave Cummings, Jeremiah Currier, Yunxing Dai, Cory Decareaux, Thomas Degry, Noah Deutsch,
442 Damien Deville, Arka Dhar, David Dohan, Steve Dowling, Sheila Dunning, Adrien Ecoffet, Atty
443 Eleti, Tyna Eloundou, David Farhi, Liam Fedus, Niko Felix, Simón Posada Fishman, Juston Forte,
444 Isabella Fulford, Leo Gao, Elie Georges, Christian Gibson, Vik Goel, Tarun Gogineni, Gabriel
445 Goh, Rapha Gontijo-Lopes, Jonathan Gordon, Morgan Grafstein, Scott Gray, Ryan Greene, Joshua
446 Gross, Shixiang Shane Gu, Yufei Guo, Chris Hallacy, Jesse Han, Jeff Harris, Yuchen He, Mike
447 Heaton, Johannes Heidecke, Chris Hesse, Alan Hickey, Wade Hickey, Peter Hoeschele, Brandon
448 Houghton, Kenny Hsu, Shengli Hu, Xin Hu, Joost Huizinga, Shantanu Jain, Shawn Jain, Joanne
449 Jang, Angela Jiang, Roger Jiang, Haozhun Jin, Denny Jin, Shino Jomoto, Billie Jonn, Heewoo
450 Jun, Tomer Kaftan, Łukasz Kaiser, Ali Kamali, Ingmar Kanitscheider, Nitish Shirish Keskar,
451 Tabarak Khan, Logan Kilpatrick, Jong Wook Kim, Christina Kim, Yongjik Kim, Jan Hendrik
452 Kirchner, Jamie Kiros, Matt Knight, Daniel Kokotajlo, Łukasz Kondraciuk, Andrew Kondrich,

453 Aris Konstantinidis, Kyle Kopic, Gretchen Krueger, Vishal Kuo, Michael Lampe, Ikai Lan, Teddy
454 Lee, Jan Leike, Jade Leung, Daniel Levy, Chak Ming Li, Rachel Lim, Molly Lin, Stephanie
455 Lin, Mateusz Litwin, Theresa Lopez, Ryan Lowe, Patricia Lue, Anna Makanju, Kim Malfacini,
456 Sam Manning, Todor Markov, Yaniv Markovski, Bianca Martin, Katie Mayer, Andrew Mayne,
457 Bob McGrew, Scott Mayer McKinney, Christine McLeavey, Paul McMillan, Jake McNeil, David
458 Medina, Aalok Mehta, Jacob Menick, Luke Metz, Andrey Mishchenko, Pamela Mishkin, Vinnie
459 Monaco, Evan Morikawa, Daniel Mossing, Tong Mu, Mira Murati, Oleg Murk, David Mély,
460 Ashvin Nair, Reiichiro Nakano, Rajeev Nayak, Arvind Neelakantan, Richard Ngo, Hyeonwoo
461 Noh, Long Ouyang, Cullen O’Keefe, Jakub Pachocki, Alex Paino, Joe Palermo, Ashley Pantuliano,
462 Giambattista Parascandolo, Joel Parish, Emy Parparita, Alex Passos, Mikhail Pavlov, Andrew Peng,
463 Adam Perelman, Filipe de Avila Belbute Peres, Michael Petrov, Henrique Ponde de Oliveira Pinto,
464 Michael, Pokorný, Michelle Pokrass, Vitchyr H. Pong, Tolly Powell, Alethea Power, Boris Power,
465 Elizabeth Proehl, Raul Puri, Alec Radford, Jack Rae, Aditya Ramesh, Cameron Raymond, Francis
466 Real, Kendra Rimbach, Carl Ross, Bob Rotsted, Henri Roussez, Nick Ryder, Mario Saltarelli, Ted
467 Sanders, Shibani Santurkar, Girish Sastry, Heather Schmidt, David Schnurr, John Schulman, Daniel
468 Selsam, Kyla Sheppard, Toki Sherbakov, Jessica Shieh, Sarah Shoker, Pranav Shyam, Szymon
469 Sidor, Eric Sigler, Maddie Simens, Jordan Sitkin, Katarina Slama, Ian Sohl, Benjamin Sokolowsky,
470 Yang Song, Natalie Staudacher, Felipe Petroski Such, Natalie Summers, Ilya Sutskever, Jie
471 Tang, Nikolas Tezak, Madeleine B. Thompson, Phil Tillet, Amin Tootoonchian, Elizabeth Tseng,
472 Preston Tuggle, Nick Turley, Jerry Tworek, Juan Felipe Cerón Uribe, Andrea Vallone, Arun
473 Vijayvergiya, Chelsea Voss, Carroll Wainwright, Justin Jay Wang, Alvin Wang, Ben Wang,
474 Jonathan Ward, Jason Wei, CJ Weinmann, Akila Welihinda, Peter Welinder, Jiayi Weng, Lilian
475 Weng, Matt Wiethoff, Dave Willner, Clemens Winter, Samuel Wolrich, Hannah Wong, Lauren
476 Workman, Sherwin Wu, Jeff Wu, Michael Wu, Kai Xiao, Tao Xu, Sarah Yoo, Kevin Yu, Qiming
477 Yuan, Wojciech Zaremba, Rowan Zellers, Chong Zhang, Marvin Zhang, Shengjia Zhao, Tianhao
478 Zheng, Juntang Zhuang, William Zhuk, and Barret Zoph. Gpt-4 technical report, 2024. URL
479 <https://arxiv.org/abs/2303.08774>.

480 Colin Rivera, Xinyi Ye, Yonsei Kim, and Wenpeng Li. Linguistic assertiveness affects factuality
481 ratings and model behavior in qa systems. In *Findings of the Association for Computational*
482 *Linguistics (ACL)*, 2023. URL <https://arxiv.org/abs/2305.04745>.

483 Siyuan Song, Jennifer Hu, and Kyle Mahowald. Language models fail to introspect about their
484 knowledge of language. *arXiv preprint arXiv:2503.07513*, 2025. URL <https://arxiv.org/abs/2503.07513>.

486 Katherine Tian, Eric Mitchell, Allan Zhou, Archit Sharma, Rafael Rafailov, Huaxiu Yao, Chelsea
487 Finn, and Christopher D. Manning. Just ask for calibration: Strategies for eliciting calibrated
488 confidence scores from language models fine-tuned with human feedback. In *Proceedings of the*
489 *2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2023. URL
490 <https://arxiv.org/abs/2305.14975>.

491 Bingbing Wen, Chenjun Xu, Bin HAN, Robert Wolfe, Lucy Lu Wang, and Bill Howe. From human
492 to model overconfidence: Evaluating confidence dynamics in large language models. In *NeurIPS*
493 *2024 Workshop on Behavioral Machine Learning*, 2024. URL [https://openreview.net/](https://openreview.net/forum?id=y9Ud05cmHs)
494 [forum?id=y9Ud05cmHs](https://openreview.net/forum?id=y9Ud05cmHs).

495 Peter West and Christopher Potts. Base models beat aligned models at randomness and creativity,
496 2025. URL <https://arxiv.org/abs/2505.00047>.

497 Bryan Wilie, Samuel Cahyawijaya, Etsuko Ishii, Junxian He, and Pascale Fung. Belief revision: The
498 adaptability of large language models reasoning, 2024. URL [https://arxiv.org/abs/2406.](https://arxiv.org/abs/2406.19764)
499 [19764](https://arxiv.org/abs/2406.19764).

500 Miao Xiong, Zhiyuan Hu, Xinyang Lu, Yifei Li, Jie Fu, Junxian He, and Bryan Hooi. Can llms
501 express their uncertainty? an empirical evaluation of confidence elicitation in llms. In *Proceedings*
502 *of the 2024 International Conference on Learning Representations (ICLR)*, 2024. URL [https:](https://arxiv.org/abs/2306.13063)
503 [//arxiv.org/abs/2306.13063](https://arxiv.org/abs/2306.13063).

504 Rongwu Xu, Brian S. Lin, Han Qiu, et al. The earth is flat because...: Investigating llms’ belief
505 towards misinformation via persuasive conversation. *arXiv preprint arXiv:2312.06717*, 2023. URL
506 <https://arxiv.org/abs/2312.06717>.

507 Yuxiang Zheng, Dayuan Fu, Xiangkun Hu, Xiaojie Cai, Lyumanshan Ye, Pengrui Lu, and Pengfei
508 Liu. Deepresearcher: Scaling deep research via reinforcement learning in real-world environments,
509 2025. URL <https://arxiv.org/abs/2504.03160>.

510 Kaitlyn Zhou, Dan Jurafsky, and Tatsunori Hashimoto. Navigating the grey area: How expressions of
511 uncertainty and overconfidence affect language models. In *Proceedings of the 2023 Conference on*
512 *Empirical Methods in Natural Language Processing (EMNLP)*, 2023a. URL <https://arxiv.org/abs/2302.13439>.

514 Kaitlyn Zhou, Dan Jurafsky, and Tatsunori Hashimoto. Navigating the grey area: How expressions of
515 uncertainty and overconfidence affect language models, 2023b. URL <https://arxiv.org/abs/2302.13439>.

517 A LLMs in the Debater Pool

518 All experiments were performed between February and May 2025

Provider	Model
openai	o3-mini
google	gemini-2.0-flash-001
anthropic	claude-3.7-sonnet
deepseek	deepseek-chat
519 qwen	qwq-32b
openai	gpt-4o-mini
google	gemma-3-27b-it
anthropic	claude-3.5-haiku
deepseek	deepseek-r1-distill-qwen-14b
qwen	qwen-max

520 B Debate Pairings Schedule

521 The debate pairings for this study were designed to ensure balanced experimental conditions while
522 maximizing informative comparisons. We employed a two-phase pairing strategy that combined
523 structured assignments with performance-based matching.

524 B.1 Pairing Objectives and Constraints

525 Our pairing methodology addressed several key requirements:

- 526 • **Equal debate opportunity:** Each model participated in 10-12 debates
- 527 • **Role balance:** Models were assigned to proposition and opposition roles with approximately
528 equal frequency
- 529 • **Opponent diversity:** Models faced a variety of opponents rather than repeatedly debating
530 the same models
- 531 • **Topic variety:** Each model-pair debated different topics to avoid topic-specific advantages
- 532 • **Performance-based matching:** After initial rounds, models with similar win-loss records
533 were paired to ensure competitive matches

534 B.2 Initial Round Planning

535 The first set of debates used predetermined pairings designed to establish baseline performance
536 metrics. These initial matchups ensured each model:

- 537 • Participated in at least two debates (one as proposition, one as opposition)
- 538 • Faced opponents from different model families (e.g., ensuring OpenAI models debated
539 against non-OpenAI models)
- 540 • Was assigned to different topics to avoid topic-specific advantages

B.3 Dynamic Performance-Based Matching

For subsequent rounds, we implemented a Swiss-tournament-style system where models were paired based on their current win-loss records and confidence calibration metrics. This approach:

1. Ranked models by performance (primary: win-loss differential, secondary: confidence margin)
2. Grouped models with similar performance records
3. Generated pairings within these groups, avoiding rematches where possible
4. Ensured balanced proposition/opposition role assignments

When an odd number of models existed in a performance tier, one model was paired with a model from an adjacent tier, prioritizing models that had not previously faced each other.

B.4 Rebalancing Rounds

After the dynamic rounds, we conducted a final set of rebalancing debates using the algorithm described in the main text. This phase ensured that any remaining imbalances in participation or role assignment were addressed, guaranteeing methodological consistency across the dataset.

Table 4: Model Debate Participation Distribution

Model	Proposition	Opposition	Total
google/gemma-3-27b-it	6	6	12
google/gemini-2.0-flash-001	6	6	12
qwen/qwen-max	6	6	12
anthropic/claude-3.5-haiku	6	6	12
qwen/qwq-32b:free	6	6	12
anthropic/claude-3.7-sonnet	6	7	13
deepseek/deepseek-chat	6	6	12
openai/gpt-4o-mini	6	6	12
openai/o3-mini	6	6	12
deepseek/deepseek-r1-distill-qwen-14b:free	6	5	11
Total debates	60	60	120

As shown in the table, the pairing schedule achieved nearly perfect balance, with eight models participating in exactly 12 debates (6 as proposition and 6 as opposition). Only two models (openai/gpt-4o-mini and deepseek/deepseek-r1-distill-qwen-14b) had slight imbalances with 11 total debates each.

This balanced design ensured that observed confidence patterns were not artifacts of pairing methodology but rather reflected genuine metacognitive properties of the models being studied.

C Debater Prompt Structures

C.1 Opening Speech

OPENING SPEECH STRUCTURE

ARGUMENT 1

Core Claim: (State your first main claim in one clear sentence)

Support Type: (Choose either EVIDENCE or PRINCIPLE)

Support Details:

For Evidence:

573 - Provide specific examples with dates/numbers
574 - Include real world cases and outcomes
575 - Show clear relevance to the topic
576 For Principle:
577 - Explain the key principle/framework
578 - Show why it is valid/important
579 - Demonstrate how it applies here
580 Connection: (Explicit explanation of how this evidence/principle proves your claim)
581
582 ARGUMENT 2
583 (Use exact same structure as Argument 1)
584
585 ARGUMENT 3 (Optional)
586 (Use exact same structure as Argument 1)
587
588 SYNTHESIS
589 - Explain how your arguments work together as a unified case
590 - Show why these arguments prove your side of the motion
591 - Present clear real-world impact and importance
592 - Link back to key themes/principles
593
594 - Follow structure exactly as shown
595 - Keep all section headers
596 - Fill in all components fully
597 - Be specific and detailed
598 - Use clear organization
599 - Label all sections
600 - No skipping components
601 JUDGING GUIDANCE
602
603 The judge will evaluate your speech using these strict criteria:
604
605 DIRECT CLASH ANALYSIS
606 - Every disagreement must be explicitly quoted and directly addressed
607 - Simply making new arguments without engaging opponents' points will be penalized
608 - Show exactly how your evidence/reasoning defeats theirs
609 - Track and reference how arguments evolve through the debate
610
611 EVIDENCE QUALITY HIERARCHY
612 1. Strongest: Specific statistics, named examples, verifiable cases with dates/numbers
613 2. Medium: Expert testimony with clear sourcing
614 3. Weak: General examples, unnamed cases, theoretical claims without support
615 - Correlation vs. causation will be scrutinized - prove causal links
616 - Evidence must directly support the specific claim being made
617
618 LOGICAL VALIDITY
619 - Each argument requires explicit warrants (reasons why it's true)
620 - All logical steps must be clearly shown, not assumed
621 - Internal contradictions severely damage your case
622 - Hidden assumptions will be questioned if not defended
623
624 RESPONSE OBLIGATIONS
625 - Every major opposing argument must be addressed
626 - Dropped arguments are considered conceded
627 - Late responses (in final speech) to early arguments are discounted
628 - Shifting or contradicting your own arguments damages credibility
629
630 IMPACT ANALYSIS & WEIGHING
631 - Explain why your arguments matter more than opponents'

- 632 - Compare competing impacts explicitly
- 633 - Show both philosophical principles and practical consequences
- 634 - Demonstrate how winning key points proves the overall motion
- 635
- 636 The judge will ignore speaking style, rhetoric, and presentation. Focus entirely on argument
- 637

638 C.2 Rebuttal Speech

639

640

641 REBUTTAL STRUCTURE

642

643 CLASH POINT 1

644 Original Claim: (Quote opponent's exact claim you're responding to)

645 Challenge Type: (Choose one)

- 646 - Evidence Critique (showing flaws in their evidence)
- 647 - Principle Critique (showing limits of their principle)
- 648 - Counter Evidence (presenting stronger opposing evidence)
- 649 - Counter Principle (presenting superior competing principle)

650 Challenge:

651 For Evidence Critique:

- 652 - Identify specific flaws/gaps in their evidence
- 653 - Show why the evidence doesn't prove their point
- 654 - Provide analysis of why it's insufficient

655 For Principle Critique:

- 656 - Show key limitations of their principle
- 657 - Demonstrate why it doesn't apply well here
- 658 - Explain fundamental flaws in their framework

659 For Counter Evidence:

- 660 - Present stronger evidence that opposes their claim
- 661 - Show why your evidence is more relevant/compelling
- 662 - Directly compare strength of competing evidence

663 For Counter Principle:

- 664 - Present your competing principle/framework
- 665 - Show why yours is superior for this debate
- 666 - Demonstrate better application to the topic

667 Impact: (Explain exactly why winning this point is crucial for the debate)

668

669 CLASH POINT 2

670 (Use exact same structure as Clash Point 1)

671

672 CLASH POINT 3

673 (Use exact same structure as Clash Point 1)

674

675 DEFENSIVE ANALYSIS

676 Vulnerabilities:

- 677 - List potential weak points in your responses
- 678 - Identify areas opponent may attack
- 679 - Show awareness of counter-arguments

680 Additional Support:

- 681 - Provide reinforcing evidence/principles
- 682 - Address likely opposition responses
- 683 - Strengthen key claims

684 Why We Prevail:

- 685 - Clear comparison of competing arguments
- 686 - Show why your responses are stronger
- 687 - Link to broader debate themes

688

689 WEIGHING
690 Key Clash Points:
691 - Identify most important disagreements
692 - Show which points matter most and why
693 Why We Win:
694 - Explain victory on key points
695 - Compare strength of competing claims
696 Overall Impact:
697 - Show how winning key points proves case
698 - Demonstrate importance for motion
699
700 - Follow structure exactly as shown
701 - Keep all section headers
702 - Fill in all components fully
703 - Be specific and detailed
704 - Use clear organization
705 - Label all sections
706 - No skipping components
707
708 JUDGING GUIDANCE
709
710 The judge will evaluate your speech using these strict criteria:
711
712 DIRECT CLASH ANALYSIS
713 - Every disagreement must be explicitly quoted and directly addressed
714 - Simply making new arguments without engaging opponents' points will be penalized
715 - Show exactly how your evidence/reasoning defeats theirs
716 - Track and reference how arguments evolve through the debate
717
718 EVIDENCE QUALITY HIERARCHY
719 1. Strongest: Specific statistics, named examples, verifiable cases with dates/numbers
720 2. Medium: Expert testimony with clear sourcing
721 3. Weak: General examples, unnamed cases, theoretical claims without support
722 - Correlation vs. causation will be scrutinized - prove causal links
723 - Evidence must directly support the specific claim being made
724
725 LOGICAL VALIDITY
726 - Each argument requires explicit warrants (reasons why it's true)
727 - All logical steps must be clearly shown, not assumed
728 - Internal contradictions severely damage your case
729 - Hidden assumptions will be questioned if not defended
730
731 RESPONSE OBLIGATIONS
732 - Every major opposing argument must be addressed
733 - Dropped arguments are considered conceded
734 - Late responses (in final speech) to early arguments are discounted
735 - Shifting or contradicting your own arguments damages credibility
736
737 IMPACT ANALYSIS & WEIGHING
738 - Explain why your arguments matter more than opponents'
739 - Compare competing impacts explicitly
740 - Show both philosophical principles and practical consequences
741 - Demonstrate how winning key points proves the overall motion
742
743 The judge will ignore speaking style, rhetoric, and presentation. Focus entirely on argument
744
745

746 C.3 Closing Speech

747

748

749

750 FINAL SPEECH STRUCTURE

751

752 FRAMING

753 Core Questions:

- 754 - Identify fundamental issues in debate
- 755 - Show what key decisions matter
- 756 - Frame how debate should be evaluated

757

758 KEY CLASHES

759 For each major clash:

760 Quote: (Exact disagreement between sides)

761 Our Case Strength:

- 762 - Show why our evidence/principles are stronger
- 763 - Provide direct comparison of competing claims
- 764 - Demonstrate superior reasoning/warrants

765 Their Response Gaps:

- 766 - Identify specific flaws in opponent response
- 767 - Show what they failed to address
- 768 - Expose key weaknesses

769 Crucial Impact:

- 770 - Explain why this clash matters
- 771 - Show importance for overall motion
- 772 - Link to core themes/principles

773

774 VOTING ISSUES

775 Priority Analysis:

- 776 - Identify which clashes matter most
- 777 - Show relative importance of points
- 778 - Clear weighing framework

779 Case Proof:

- 780 - How winning key points proves our case
- 781 - Link arguments to motion
- 782 - Show logical chain of reasoning

783 Final Weighing:

- 784 - Why any losses don't undermine case
- 785 - Overall importance of our wins
- 786 - Clear reason for voting our side

787

- 788 - Follow structure exactly as shown
- 789 - Keep all section headers
- 790 - Fill in all components fully
- 791 - Be specific and detailed
- 792 - Use clear organization
- 793 - Label all sections
- 794 - No skipping components

795

796 JUDGING GUIDANCE

797

798 The judge will evaluate your speech using these strict criteria:

799

800 DIRECT CLASH ANALYSIS

- 801 - Every disagreement must be explicitly quoted and directly addressed
- 802 - Simply making new arguments without engaging opponents' points will be penalized
- 803 - Show exactly how your evidence/reasoning defeats theirs

804 - Track and reference how arguments evolve through the debate
805
806 EVIDENCE QUALITY HIERARCHY
807 1. Strongest: Specific statistics, named examples, verifiable cases with dates/numbers
808 2. Medium: Expert testimony with clear sourcing
809 3. Weak: General examples, unnamed cases, theoretical claims without support
810 - Correlation vs. causation will be scrutinized - prove causal links
811 - Evidence must directly support the specific claim being made
812
813 LOGICAL VALIDITY
814 - Each argument requires explicit warrants (reasons why it's true)
815 - All logical steps must be clearly shown, not assumed
816 - Internal contradictions severely damage your case
817 - Hidden assumptions will be questioned if not defended
818
819 RESPONSE OBLIGATIONS
820 - Every major opposing argument must be addressed
821 - Dropped arguments are considered conceded
822 - Late responses (in final speech) to early arguments are discounted
823 - Shifting or contradicting your own arguments damages credibility
824
825 IMPACT ANALYSIS & WEIGHING
826 - Explain why your arguments matter more than opponents'
827 - Compare competing impacts explicitly
828 - Show both philosophical principles and practical consequences
829 - Demonstrate how winning key points proves the overall motion
830
831 The judge will ignore speaking style, rhetoric, and presentation. Focus entirely on argument
832
833

834 **D AI Jury Prompt Details**

835 **D.1 Jury Selection and Validation Process**

836 Before conducting the full experiment, we performed a validation study using a set of six sample
837 debates. These validation debates were evaluated by multiple candidate judge models to assess their
838 reliability, calibration, and analytical consistency. The validation process revealed that:

- 839 • Models exhibited varying levels of agreement with human expert evaluations
- 840 • Some models showed consistent biases toward either proposition or opposition sides
- 841 • Certain models demonstrated superior ability to identify key clash points and evaluate
842 evidence quality
- 843 • Using a panel of judges rather than a single model significantly improved evaluation reliabil-
844 ity

845 Based on these findings, we selected our final jury composition of six judges: two instances each of
846 qwen/qwq-32b, google/gemini-pro-1.5, and deepseek/deepseek-chat. This combination
847 provided both architectural diversity and strong analytical performance.

848 **D.2 Jury Evaluation Protocol**

849 Each debate was independently evaluated by all six judges following this protocol:

- 850 1. Judges received the complete debate transcript with all confidence bet information removed
- 851 2. Each judge analyzed the transcript according to the criteria specified in the prompt below
- 852 3. Judges provided a structured verdict including winner determination, confidence level, and
853 detailed reasoning

854 4. The six individual judgments were aggregated to determine the final winner, with the side
855 receiving the higher sum of confidence scores declared victorious

856 D.3 Complete Judge Prompt

857 The following is the verbatim prompt provided to each AI judge:

```
858
859
860 You are an expert debate judge. Your role is to analyze formal debates using the
861   ↳ following strictly prioritized criteria:
862 I. Core Judging Principles (In order of importance):
863 Direct Clash Resolution:
864 Identify all major points of disagreement (clashes) between the teams.
865 For each clash:
866 Quote the exact statements representing each side's position.
867 Analyze the logical validity of each argument within the clash. Is the reasoning
868   ↳ sound, or does it contain fallacies (e.g., hasty generalization,
869   ↳ correlation/causation, straw man, etc.)? Identify any fallacies by name.
870 Analyze the quality of evidence presented within that specific clash. Define "
871   ↳ quality" as:
872 Direct Relevance: How directly does the evidence support the claim being made?
873   ↳ Does it establish a causal link, or merely a correlation? Explain the
874   ↳ difference if a causal link is claimed but not proven.
875 Specificity: Is the evidence specific and verifiable (e.g., statistics, named
876   ↳ examples, expert testimony), or vague and general? Prioritize specific
877   ↳ evidence.
878 Source Credibility (If Applicable): If a source is cited, is it generally
879   ↳ considered reliable and unbiased? If not, explain why this weakens the
880   ↳ evidence.
881 Evaluate the effectiveness of each side's rebuttals within the clash. Define "
882   ↳ effectiveness" as:
883 Direct Response: Does the rebuttal directly address the opponent's claim and
884   ↳ evidence? If not, explain how this weakens the rebuttal.
885 Undermining: Does the rebuttal successfully weaken the opponent's argument (e.g.,
886   ↳ by exposing flaws in logic, questioning evidence, presenting counter-
887   ↳ evidence)? Explain how the undermining occurs.
888 Explicitly state which side wins the clash and why, referencing your analysis of
889   ↳ logic, evidence, and rebuttals. Provide at least two sentences of
890   ↳ justification for each clash decision, explaining the relative strength of
891   ↳ the arguments.
892 Track the evolution of arguments through the debate within each clash. How did the
893   ↳ claims and responses change over time? Note any significant shifts or
894   ↳ concessions.
895 Argument Hierarchy and Impact:
896 Identify the core arguments of each side (the foundational claims upon which their
897   ↳ entire case rests).
898 Explain the logical links between each core argument and its supporting claims/
899   ↳ evidence. Are the links clear, direct, and strong? If not, explain why this
900   ↳ weakens the argument.
901 Assess the stated or clearly implied impacts of each argument. What are the
902   ↳ consequences if the argument is true? Be specific.
903 Determine the relative importance of each core argument to the overall debate.
904   ↳ Which arguments are most central to resolving the motion? State this
905   ↳ explicitly and justify your ranking.
906 Weighing Principled vs. Practical Arguments: When weighing principled arguments (
907   ↳ based on abstract concepts like rights or justice) against practical
908   ↳ arguments (based on real-world consequences), consider:
909 (a) the strength and universality of the underlying principle;
910 (b) the directness, strength, and specificity of the evidence supporting the
911   ↳ practical claims; and
912 (c) the extent to which the practical arguments directly address, mitigate, or
913   ↳ outweigh the concerns raised by the principled arguments. Explain your
914   ↳ reasoning.
915 Consistency and Contradictions:
```

916 Identify any internal contradictions within each team's case (arguments that
917 ↳ contradict each other).

918 Identify any inconsistencies between a team's arguments and their rebuttals.
919 Note any dropped arguments (claims made but not responded to). For each dropped
920 ↳ argument:

921 Assess its initial strength based on its logical validity and supporting evidence,
922 ↳ as if it had not been dropped.

923 Then, consider the impact of it being unaddressed. Does the lack of response
924 ↳ significantly weaken the overall case of the side that dropped it? Explain
925 ↳ why or why not.

926 II. Evaluation Requirements:

927 Steelmanning: When analyzing arguments, present them in their strongest possible
928 ↳ form, even if you disagree with them. Actively look for the most charitable
929 ↳ interpretation.

930 Argument-Based Decision: Base your decision solely on the arguments made within
931 ↳ the debate text provided. Do not introduce outside knowledge or opinions.
932 ↳ If an argument relies on an unstated assumption, analyze it only if that
933 ↳ assumption is clearly and necessarily implied by the presented arguments.

934 Ignore Presentation: Disregard presentation style, speaking quality, rhetorical
935 ↳ flourishes, etc. Focus exclusively on the substance of the arguments and
936 ↳ their logical connections.

937 Framework Neutrality: If both sides present valid but competing frameworks for
938 ↳ evaluating the debate, maintain neutrality between them. Judge the debate
939 ↳ based on how well each side argues within their chosen framework, and
940 ↳ according to the prioritized criteria in Section I.

941 III. Common Judging Errors to AVOID:

942 Intervention: Do not introduce your own arguments or evidence.

943 Shifting the Burden of Proof: Do not place a higher burden of proof on one side
944 ↳ than the other. Both sides must prove their claims to the same standard.

945 Over-reliance on "Real-World" Arguments: Do not automatically favor arguments
946 ↳ based on "real-world" examples over principled or theoretical arguments.
947 ↳ Evaluate all arguments based on the criteria in Section I.

948 Ignoring Dropped Arguments: Address all dropped arguments as specified in I.3.

949 Double-Counting: Do not give credit for the same argument multiple times.

950 Assuming Causation from Correlation: Be highly skeptical of arguments that claim
951 ↳ causation based solely on correlation. Demand clear evidence of a causal
952 ↳ mechanism.

953 Not Justifying Clash Decisions: Provide explicit justification for every clash
954 ↳ decision, as required in I.1.

955 IV. Decision Making:

956 Winner: The winner must be either "Proposition" or "Opposition" (no ties).

957 Confidence Level: Assign a confidence level (0-100) reflecting the margin of
958 ↳ victory. A score near 50 indicates a very close debate.

959 90-100: Decisive Victory
960 70-89: Clear Victory
961 51-69: Narrow Victory.

962 Explain why you assigned the specific confidence level.

963 Key Factors: Identify the 2-3 most crucial factors that determined the outcome.
964 ↳ These should be specific clashes or arguments that had the greatest impact
965 ↳ on your decision. Explain why these factors were decisive.

966 Detailed Reasoning: Provide a clear, logical, and detailed explanation for your
967 ↳ conclusion. Explain how the key factors interacted to produce the result.
968 ↳ Reference specific arguments and analysis from sections I-III. Show your
969 ↳ work, step-by-step. Do not simply state your conclusion; justify it with
970 ↳ reference to the specific arguments made.

971 V. Line-by-Line Justification:

972 Create a section titled "V. Line-by-Line Justification."

973 In this section, provide at least one sentence referencing each and every section
974 ↳ of the provided debate text (Prop 1, Opp 1, Prop Rebuttal 1, Opp Rebuttal
975 ↳ 1, Prop Final, Opp Final). This ensures that no argument, however minor,
976 ↳ goes unaddressed. You may group multiple minor arguments together in a
977 ↳ single sentence if they are closely related. The purpose is to demonstrate
978 ↳ that you have considered the entirety of the debate.

979 VI. Format for your response:

```

980 Organize your response in clearly marked sections exactly corresponding to the
981     ↪ sections above (I.1, I.2, I.3, II, III, IV, V). This structured output is
982     ↪ mandatory. Your response must follow this format to be accepted.
983
984
985
986 format:
987 write all your thoughts out
988 then put in XML tags
989 <winnerName>opposition|proposition</winnerName>
990
991 <confidence>0-100</confidence>\n
992
993 These existing is compulsory as the parser will fail otherwise
994

```

995 D.4 Evaluation Methodology: The AI Jury

996 Evaluating 60 debates rigorously required a scalable and consistent approach. We implemented an AI
997 jury system to ensure robust assessment based on argumentative merit.

998 **Rationale for AI Jury:** This approach was chosen over single AI judges (to mitigate potential bias
999 and improve reliability through aggregation) and human judges (due to the scale and cost required for
1000 consistent evaluation of this many debates).

1001 **Jury Selection Process:** Potential judge models were evaluated based on criteria including: (1) Per-
1002 formance Reliability (agreement with consensus, confidence calibration, consistency across debates),
1003 (2) Analytical Quality (ability to identify clash, evaluate evidence, recognize fallacies), (3) Diversity
1004 (representation from different model architectures and providers), and (4) Cost-Effectiveness.

1005 **Final Jury Composition:** The final jury consisted of six judges in total, comprising two instances
1006 each of qwen/qwq-32b, google/gemini-pro-1.5, and deepseek/deepseek-chat. This combi-
1007 nation provided architectural diversity from three providers, included models demonstrating strong
1008 analytical performance and calibration during selection, and balanced quality with cost. Each debate
1009 was judged independently by all six judges.

1010 **Judging Procedure & Prompt:** Judges evaluated the full debate transcript based solely on the
1011 argumentative substance presented, adhering to a highly detailed prompt (see Appendix D for full
1012 text). Key requirements included:

- 1013 • Strict focus on **Direct Clash Resolution:** Identifying, quoting, and analyzing each point
1014 of disagreement based on logic, evidence quality (using a defined hierarchy), and rebuttal
1015 effectiveness, explicitly determining a winner for each clash with justification.
- 1016 • Evaluation of **Argument Hierarchy & Impact** and overall case **Consistency**.
- 1017 • Explicit instructions to **ignore presentation style** and avoid common judging errors (e.g.,
1018 intervention, shifting burdens).
- 1019 • Requirement for **Structured Output:** Including Winner (Proposition/Opposition), Confi-
1020 dence (0-100, representing margin of victory), Key Deciding Factors, Detailed Step-by-Step
1021 Reasoning, and a **Line-by-Line Justification** section confirming review of the entire tran-
1022 script.

1023 **Final Verdict Determination:** The final winner for each debate was determined by aggregating
1024 the outputs of the six judges. The side (Proposition or Opposition) that received the higher sum of
1025 confidence scores across all six judges was declared the winner. The normalized difference between
1026 the winner’s total confidence and the loser’s total confidence served as the margin of victory. Ties in
1027 total confidence were broken randomly.

1028 E Topics of Debate

- 1029 • This House would require national television news broadcasters with over 5% annual view-
1030 ership to provide equal prime-time coverage to parties polling above 10% and guaranteed

```

===== JUDGE PROMPT (CORE EXCERPT) =====

I. CORE JUDGING PRINCIPLES
1. Direct Clash Resolution
  - Quote each disagreement
  - Analyse logic, evidence quality, rebuttal success
  - Declare winner of the clash with rationale
2. Argument Hierarchy & Impact
  - Identify each side's core arguments
  - Trace logical links and stated impacts
  - Rank which arguments decide the motion
3. Consistency & Contradictions
  - Flag internal contradictions, dropped points

II. EVALUATION REQUIREMENTS
  - Steelman arguments
  - Do NOT add outside knowledge
  - Ignore presentation style

III. COMMON JUDGING ERRORS TO AVOID
Intervention - Burden-shifting - Double-counting -
Assuming causation from correlation - Ignoring dropped arguments

IV. DECISION FORMAT
<winnerName> Proposition|Opposition </winnerName>
<confidence> 0-100 </confidence>
Key factors (2-3 bullet list)
Detailed section-by-section reasoning

V. LINE-BY-LINE JUSTIFICATION
Provide > 1 sentence addressing Prop 1, Opp 1, Rebuttals, Finals
=====

```

Figure 2: Condensed version of the judge prompt given to the AI jury (full text in Appendix D).

1031 response segments within 48 hours of criticism, rather than relying on media watchdog
1032 guidelines and voluntary fairness codes

- 1033 • This House would require US state governors to face recall elections through voter petitions
1034 (requiring 20% of registered voters within 90 days) rather than allowing removal during
1035 their term only through state legislative impeachment, with both mechanisms prohibited
1036 during the first and last 6 months of their term
- 1037 • This House believes that governments should transition their primary role in space from
1038 direct operation to regulation and oversight of private sector space activities
- 1039 • This House believes that professors should actively engage in public advocacy on social and
1040 political issues within their field of expertise
- 1041 • This House would require G20 nations to participate in a unified carbon trading market
1042 with cross-border credit trading and quarterly auctions, rather than allowing each nation to
1043 implement its own domestic carbon tax system
- 1044 • This House would limit individual shareholding in social media platforms with over 100 mil-
1045 lion monthly active users to a maximum of 15% voting rights, requiring broader institutional
1046 and public ownership instead of allowing concentrated private control

1047 **F Self Debate Ablation**

1048 We appended this text to the prompt given to the LLM:"NOTE: In this debate, you are
1049 debating against yourself. The other debater is without question another
1050 copy of yourself and has the same competence as you have"

1051 **G Informed Self Debate Ablation**

1052 We appended this text to the prompt given to the LLM:"NOTE: In this debate, you are
1053 debating against yourself. The other debater is without question another
1054 copy of yourself and has the same competence as you have. By definition,

1055 you have exactly a 50% chance of winning this debate since you are debating
1056 an identical copy of yourself"

1057 H Public Self Debate Ablation

1058 We appended this text to the prompt given to the LLM:"NOTE: In this debate, you are
1059 debating against yourself. The other debater is without question another
1060 copy of yourself and has the same competence as you have. By definition,
1061 you have exactly a 50% chance of winning this debate since you are debating
1062 an identical copy of yourself. ...

1063 After your speech, you must include a public confidence bet (0-100)
1064 indicating how likely you think you are to win this debate"

1065 I Hypothesis Tests

1066 **Test for General Overconfidence in Opening Statements** To statistically evaluate the hypothesis
1067 that LLMs exhibit general overconfidence in their initial self-assessments, we performed a one-sample
1068 t-test. This test compares the mean of a sample to a known or hypothesized population mean. The data
1069 used for this test was the collection of all opening confidence bets submitted by both Proposition and
1070 Opposition debaters across all 60 debates (total N=120 individual opening bets). The null hypothesis
1071 (H_0) was that the mean of these opening confidence bets was equal to 50% (the expected win rate in
1072 a fair, symmetric contest). The alternative hypothesis (H_1) was that the mean was greater than 50%,
1073 reflecting pervasive overconfidence. The analysis yielded a mean opening confidence of 72.92%.
1074 The results of the one-sample t-test were $t = 31.666$, with a one-tailed $p < 0.0001$. With a p-value
1075 well below the standard significance level of 0.05, we reject the null hypothesis. This provides
1076 strong statistical evidence that the average opening confidence level of LLMs in this debate setting is
1077 significantly greater than the expected 50%, supporting the claim of pervasive initial overconfidence.

1078 J Detailed Initial Confidence Test Results

1079 This appendix provides the full results of the one-sample hypothesis tests conducted for the mean
1080 initial confidence of each language model within each experimental configuration. The tests assess
1081 whether the mean reported confidence is statistically significantly greater than 50%.

1082 K Detailed Confidence Escalation Results

1083 This appendix provides the full details of the confidence escalation analysis across rounds (Opening,
1084 Rebuttal, Closing) for each language model within each experimental configuration. We analyze the
1085 change in mean confidence between rounds using paired statistical tests to assess the significance of
1086 escalation.

1087 For each experiment type and model, we report the mean confidence (\pm Standard Deviation, N) for
1088 each round. We then report the mean difference (Δ) in confidence between rounds (Later Round
1089 Bet - Earlier Round Bet) and the p-value from a one-sided paired t-test (H_1 : Later Round Bet >
1090 Earlier Round Bet). A significant positive Δ indicates statistically significant confidence escalation
1091 during that transition. For completeness, we also include the results of two-sided Wilcoxon signed-
1092 rank tests where applicable. Significance levels are denoted as: * $p \leq 0.05$, ** $p \leq 0.01$, *** $p \leq 0.001$.

1093 Note that for transitions where there was no variance in the bet differences (e.g., all changes were
1094 exactly 0), the p-value for the t-test is indeterminate or the test is not applicable. In such cases, we
1095 indicate '-' and rely on the mean difference ($\Delta = 0.00$) and the mean values themselves (which are
1096 equal). The Wilcoxon test might also yield non-standard results or N/A in some low-variance cases.

1097 K.1 Confidence Escalation by Experiment Type and Model

Table 5: One-Sample Hypothesis Test Results for Mean Initial Confidence (vs. 50%). Tests were conducted for each model in each configuration against the null hypothesis that the true mean initial confidence is $\leq 50\%$. Significant results ($p \leq 0.05$) indicate statistically significant overconfidence. Results from both t-tests and Wilcoxon signed-rank tests are provided.

Experiment	Model	N	Mean	t-test vs 50% (H1: > 50)		Wilcoxon vs 50% (H1: > 50)	
				p-value	Significant	p-value	Significant
Cross-model	qwen/qwen-max	12	73.33	6.97×10^{-7}	True	0.0002	True
Cross-model	anthropic/claude-3.5-haiku	12	71.67	4.81×10^{-9}	True	0.0002	True
Cross-model	deepseek/deepseek-r1-distill-qwen-14b:free	11	79.09	1.64×10^{-6}	True	0.0005	True
Cross-model	anthropic/claude-3.7-sonnet	13	67.31	8.76×10^{-10}	True	0.0001	True
Cross-model	google/gemini-2.0-flash-001	12	65.42	2.64×10^{-5}	True	0.0007	True
Cross-model	qwen/qwq-32b:free	12	78.75	5.94×10^{-11}	True	0.0002	True
Cross-model	google/gemma-3-27b-it	12	67.50	4.74×10^{-7}	True	0.0002	True
Cross-model	openai/gpt-4o-mini	12	75.00	4.81×10^{-11}	True	0.0002	True
Cross-model	openai/o3-mini	12	77.50	2.34×10^{-9}	True	0.0002	True
Cross-model	deepseek/deepseek-chat	12	74.58	6.91×10^{-8}	True	0.0002	True
Debate against same model	qwen/qwen-max	12	62.08	0.0039	True	0.0093	True
Debate against same model	anthropic/claude-3.5-haiku	12	71.25	9.58×10^{-8}	True	0.0002	True
Debate against same model	deepseek/deepseek-r1-distill-qwen-14b:free	12	76.67	1.14×10^{-5}	True	0.0002	True
Debate against same model	anthropic/claude-3.7-sonnet	12	56.25	0.0140	True	0.0159	True
Debate against same model	google/gemini-2.0-flash-001	12	43.25	0.7972	False	0.8174	False
Debate against same model	qwen/qwq-32b:free	12	70.83	1.49×10^{-5}	True	0.0002	True
Debate against same model	google/gemma-3-27b-it	12	68.75	1.38×10^{-6}	True	0.0002	True
Debate against same model	openai/gpt-4o-mini	12	67.08	2.58×10^{-6}	True	0.0005	True
Debate against same model	openai/o3-mini	12	70.00	2.22×10^{-5}	True	0.0005	True
Debate against same model	deepseek/deepseek-chat	12	54.58	0.0043	True	0.0156	True
Informed Self (50% informed)	qwen/qwen-max	12	43.33	0.8388	False	0.7451	False
Informed Self (50% informed)	anthropic/claude-3.5-haiku	12	54.58	0.0640	False	0.0845	False
Informed Self (50% informed)	deepseek/deepseek-r1-distill-qwen-14b:free	12	55.75	0.0007	True	0.0039	True
Informed Self (50% informed)	anthropic/claude-3.7-sonnet	12	50.08	0.4478	False	0.5000	False
Informed Self (50% informed)	google/gemini-2.0-flash-001	12	36.25	0.9527	False	0.7976	False
Informed Self (50% informed)	qwen/qwq-32b:free	12	50.42	0.1694	False	0.5000	False
Informed Self (50% informed)	google/gemma-3-27b-it	12	53.33	0.1612	False	0.0820	False
Informed Self (50% informed)	openai/gpt-4o-mini	12	57.08	0.0397	True	0.0525	False
Informed Self (50% informed)	openai/o3-mini	12	50.00	— ¹	False	— ²	False
Informed Self (50% informed)	deepseek/deepseek-chat	12	49.17	0.6712	False	0.6250	False
Public Bets	qwen/qwen-max	12	64.58	0.0004	True	0.0012	True
Public Bets	anthropic/claude-3.5-haiku	12	73.33	1.11×10^{-7}	True	0.0002	True
Public Bets	deepseek/deepseek-r1-distill-qwen-14b:free	12	69.58	0.0008	True	0.0056	True
Public Bets	anthropic/claude-3.7-sonnet	12	56.25	0.0022	True	0.0054	True
Public Bets	google/gemini-2.0-flash-001	12	34.58	0.9686	False	0.9705	False
Public Bets	qwen/qwq-32b:free	12	71.67	1.44×10^{-6}	True	0.0002	True
Public Bets	google/gemma-3-27b-it	12	63.75	0.0003	True	0.0017	True
Public Bets	openai/gpt-4o-mini	12	72.92	3.01×10^{-9}	True	0.0002	True
Public Bets	openai/o3-mini	12	72.08	2.79×10^{-6}	True	0.0002	True
Public Bets	deepseek/deepseek-chat	12	56.25	0.0070	True	0.0137	True

Table 6: Mean (\pm SD, N) Confidence and Paired Test Results for Confidence Escalation in Cross-model Debates.

Model	Opening Bet	Rebuttal Bet	Closing Bet	Open→Rebuttal	Rebuttal→Closing	Open→Closing
anthropic/claude-3.5-haiku	71.67 \pm 4.71 (N=12)	73.75 \pm 12.93 (N=12)	83.33 \pm 7.45 (N=12)	$\Delta=2.08$, p=0.2658	$\Delta=9.58$, p=0.0036**	$\Delta=11.67$, p=0.0006***
anthropic/claude-3.7-sonnet	67.31 \pm 3.73 (N=13)	73.85 \pm 4.45 (N=13)	82.69 \pm 5.04 (N=13)	$\Delta=6.54$, p=0.0003***	$\Delta=8.85$, p=0.0000***	$\Delta=15.38$, p=0.0000***
deepseek/deepseek-chat	74.58 \pm 6.91 (N=12)	77.92 \pm 9.67 (N=12)	80.00 \pm 8.66 (N=12)	$\Delta=3.33$, p=0.1099	$\Delta=2.08$, p=0.1049	$\Delta=5.42$, p=0.0077**
deepseek/deepseek-r1-distill-qwen-14b:free	79.09 \pm 9.96 (N=11)	80.45 \pm 10.76 (N=11)	86.36 \pm 9.32 (N=11)	$\Delta=1.36$, p=0.3474	$\Delta=5.91$, p=0.0172*	$\Delta=7.27$, p=0.0229*
google/gemini-2.0-flash-001	65.42 \pm 8.03 (N=12)	63.75 \pm 7.40 (N=12)	64.00 \pm 7.20 (N=12)	$\Delta=1.67$, p=0.7152	$\Delta=0.25$, p=0.4571	$\Delta=1.42$, p=0.6508
google/gemma-3-27b-it	67.50 \pm 5.95 (N=12)	78.33 \pm 5.53 (N=12)	88.33 \pm 5.14 (N=12)	$\Delta=10.83$, p=0.0000***	$\Delta=10.00$, p=0.0001***	$\Delta=20.83$, p=0.0000***
gpt-4o-mini	75.00 \pm 3.54 (N=12)	78.33 \pm 4.71 (N=12)	82.08 \pm 5.94 (N=12)	$\Delta=3.33$, p=0.0272*	$\Delta=3.75$, p=0.0008***	$\Delta=7.08$, p=0.0030**
o3-mini	77.50 \pm 5.59 (N=12)	81.25 \pm 4.15 (N=12)	84.50 \pm 3.93 (N=12)	$\Delta=3.75$, p=0.0001***	$\Delta=3.25$, p=0.0020**	$\Delta=7.00$, p=0.0001***
qwen-max	73.33 \pm 8.25 (N=12)	81.92 \pm 7.61 (N=12)	88.75 \pm 9.16 (N=12)	$\Delta=8.58$, p=0.0001***	$\Delta=6.83$, p=0.0007***	$\Delta=15.42$, p=0.0002***
qwq-32b:free	78.75 \pm 4.15 (N=12)	87.67 \pm 3.97 (N=12)	92.83 \pm 4.43 (N=12)	$\Delta=8.92$, p=0.0000***	$\Delta=5.17$, p=0.0000***	$\Delta=14.08$, p=0.0000***
OVERALL	72.92 \pm 7.89 (N=120)	77.67 \pm 9.75 (N=120)	83.26 \pm 10.06 (N=120)	$\Delta=4.75$, p<0.001***	$\Delta=5.59$, p<0.001***	$\Delta=10.34$, p<0.001***

Table 7: Mean (\pm SD, N) Confidence and Paired Test Results for Confidence Escalation in Informed Self Debates.

Model	Opening Bet	Rebuttal Bet	Closing Bet	Open→Rebuttal	Rebuttal→Closing	Open→Closing
claude-3.5-haiku	54.58 \pm 9.23 (N=12)	63.33 \pm 5.89 (N=12)	61.25 \pm 5.45 (N=12)	$\Delta=8.75$, p=0.0243*	$\Delta=2.08$, p=0.7891	$\Delta=6.67$, p=0.0194*
claude-3.7-sonnet	50.08 \pm 2.06 (N=12)	54.17 \pm 2.76 (N=12)	54.33 \pm 2.56 (N=12)	$\Delta=4.08$, p=0.0035**	$\Delta=0.17$, p=0.4190	$\Delta=4.25$, p=0.0019**
deepseek-chat	49.17 \pm 6.07 (N=12)	52.92 \pm 3.20 (N=12)	55.00 \pm 3.54 (N=12)	$\Delta=3.75$, p=0.0344*	$\Delta=2.08$, p=0.1345	$\Delta=5.83$, p=0.0075**
deepseek-r1-distill-qwen-14b:free	55.75 \pm 4.51 (N=12)	59.58 \pm 14.64 (N=12)	57.58 \pm 9.40 (N=12)	$\Delta=3.83$, p=0.1824	$\Delta=2.00$, p=0.2607	$\Delta=1.83$, p=0.2607
google/gemini-2.0-flash-001	36.25 \pm 24.93 (N=12)	50.50 \pm 11.27 (N=12)	53.92 \pm 14.53 (N=12)	$\Delta=14.25$, p=0.0697	$\Delta=3.42$, p=0.2816	$\Delta=17.67$, p=0.0211**
gemma-3-27b-it	53.33 \pm 10.67 (N=12)	57.08 \pm 10.10 (N=12)	60.83 \pm 10.96 (N=12)	$\Delta=3.75$, p=0.2279	$\Delta=3.75$, p=0.1527	$\Delta=7.50$, p=0.0859
gpt-4o-mini	57.08 \pm 12.15 (N=12)	63.75 \pm 7.67 (N=12)	65.83 \pm 8.12 (N=12)	$\Delta=6.67$, p=0.0718	$\Delta=2.08$, p=0.1588	$\Delta=8.75$, p=0.0255*
o3-mini	50.00 \pm 0.00 (N=12)	52.08 \pm 3.20 (N=12)	50.00 \pm 0.00 (N=12)	$\Delta=2.08$, p=0.0269*	$\Delta=2.08$, p=0.9731	$\Delta=0.00$, p=— ³
qwen-max	43.33 \pm 21.34 (N=12)	54.17 \pm 12.56 (N=12)	61.67 \pm 4.71 (N=12)	$\Delta=10.83$, p=0.0753	$\Delta=7.50$, p=0.0475*	$\Delta=18.33$, p=0.0124*
qwq-32b:free	50.42 \pm 1.38 (N=12)	50.08 \pm 0.28 (N=12)	50.42 \pm 1.38 (N=12)	$\Delta=0.33$, p=0.7716	$\Delta=0.33$, p=0.2284	$\Delta=0.00$, p=0.5000
OVERALL	50.00 \pm 13.55 (N=120)	55.77 \pm 9.73 (N=120)	57.08 \pm 8.97 (N=120)	$\Delta=5.77$, p<0.001***	$\Delta=1.32$, p=0.0945	$\Delta=7.08$, p<0.001***

Table 8: Mean (\pm SD, N) Confidence and Paired Test Results for Confidence Escalation in Public Bets Debates.

Model	Opening Bet	Rebuttal Bet	Closing Bet	Open→Rebuttal	Rebuttal→Closing	Open→Closing
claude-3.5-haiku	73.33 \pm 6.87 (N=12)	76.67 \pm 7.73 (N=12)	80.83 \pm 8.86 (N=12)	$\Delta=3.33$, p=0.0902	$\Delta=4.17$, p=0.0126*	$\Delta=7.50$, p=0.0117*
claude-3.7-sonnet	56.25 \pm 5.82 (N=12)	61.67 \pm 4.25 (N=12)	68.33 \pm 5.53 (N=12)	$\Delta=5.42$, p=0.0027**	$\Delta=6.67$, p=0.0016**	$\Delta=12.08$, p=0.0000***
deepseek-chat	56.25 \pm 7.11 (N=12)	62.50 \pm 6.29 (N=12)	61.67 \pm 7.73 (N=12)	$\Delta=6.25$, p=0.0032**	$\Delta=-0.83$, p=0.7247	$\Delta=5.42$, p=0.0176*
deepseek-r1-distill-qwen-14b:free	69.58 \pm 15.61 (N=12)	72.08 \pm 16.00 (N=12)	76.67 \pm 10.47 (N=12)	$\Delta=2.50$, p=0.1463	$\Delta=4.58$, p=0.0424*	$\Delta=7.08$, p=0.0136*
google/gemini-2.0-flash-001	34.58 \pm 24.70 (N=12)	44.33 \pm 21.56 (N=12)	48.25 \pm 18.88 (N=12)	$\Delta=9.75$, p=0.0195*	$\Delta=3.92$, p=0.2655	$\Delta=13.67$, p=0.0399*
gemma-3-27b-it	63.75 \pm 9.38 (N=12)	68.75 \pm 22.09 (N=12)	84.17 \pm 3.44 (N=12)	$\Delta=5.00$, p=0.2455	$\Delta=15.42$, p=0.0210*	$\Delta=20.42$, p=0.0000***
gpt-4o-mini	72.92 \pm 4.77 (N=12)	81.00 \pm 4.58 (N=12)	85.42 \pm 5.19 (N=12)	$\Delta=8.08$, p=0.0000***	$\Delta=4.42$, p=0.0004***	$\Delta=12.50$, p=0.0000***
o3-mini	72.08 \pm 9.00 (N=12)	77.92 \pm 7.20 (N=12)	80.83 \pm 6.07 (N=12)	$\Delta=5.83$, p=0.0001***	$\Delta=2.92$, p=0.0058**	$\Delta=8.75$, p=0.0001***
qwen-max	64.58 \pm 10.50 (N=12)	69.83 \pm 6.48 (N=12)	73.08 \pm 6.86 (N=12)	$\Delta=5.25$, p=0.0235*	$\Delta=3.25$, p=0.0135*	$\Delta=8.50$, p=0.0076**
qwq-32b:free	71.67 \pm 8.25 (N=12)	79.58 \pm 4.77 (N=12)	82.25 \pm 6.88 (N=12)	$\Delta=7.92$, p=0.0001***	$\Delta=2.67$, p=0.0390*	$\Delta=10.58$, p=0.0003***
OVERALL	63.50 \pm 16.31 (N=120)	69.43 \pm 16.03 (N=120)	74.15 \pm 14.34 (N=120)	$\Delta=5.93$, p<0.001***	$\Delta=4.72$, p<0.001***	$\Delta=10.65$, p<0.001***

Table 9: Mean (\pm SD, N) Confidence and Paired Test Results for Confidence Escalation in Standard Self Debates.

Model	Opening Bet	Rebuttal Bet	Closing Bet	Open→Rebuttal	Rebuttal→Closing	Open→Closing
claude-3.5-haiku	71.25 \pm 6.17 (N=12)	76.67 \pm 9.43 (N=12)	83.33 \pm 7.73 (N=12)	$\Delta=5.42$, p=0.0176*	$\Delta=6.67$, p=0.0006***	$\Delta=12.08$, p=0.0002***
claude-3.7-sonnet	56.25 \pm 8.20 (N=12)	63.33 \pm 4.25 (N=12)	68.17 \pm 6.15 (N=12)	$\Delta=7.08$, p=0.0167*	$\Delta=4.83$, p=0.0032**	$\Delta=11.92$, p=0.0047**
deepseek-chat	54.58 \pm 4.77 (N=12)	59.58 \pm 6.28 (N=12)	61.67 \pm 7.73 (N=12)	$\Delta=5.00$, p=0.0076**	$\Delta=2.08$, p=0.0876	$\Delta=7.08$, p=0.0022**
deepseek-r1-distill-qwen-14b:free	76.67 \pm 12.64 (N=12)	72.92 \pm 13.61 (N=12)	77.08 \pm 14.78 (N=12)	$\Delta=-3.75$, p=0.9591	$\Delta=4.17$, p=0.0735	$\Delta=0.42$, p=0.4570
google/gemini-2.0-flash-001	43.25 \pm 25.88 (N=12)	47.58 \pm 29.08 (N=12)	48.75 \pm 20.31 (N=12)	$\Delta=4.33$, p=0.2226	$\Delta=1.17$, p=0.4268	$\Delta=5.50$, p=0.1833
gemma-3-27b-it	68.75 \pm 7.11 (N=12)	77.92 \pm 6.60 (N=12)	85.83 \pm 6.07 (N=12)	$\Delta=9.17$, p=0.0000***	$\Delta=7.92$, p=0.0000***	$\Delta=17.08$, p=0.0000***
gpt-4o-mini	67.08 \pm 6.91 (N=12)	67.92 \pm 20.96 (N=12)	80.00 \pm 4.08 (N=12)	$\Delta=0.83$, p=0.4534	$\Delta=12.08$, p=0.0298*	$\Delta=12.92$, p=0.0002***
o3-mini	70.00 \pm 10.21 (N=12)	75.00 \pm 9.57 (N=12)	79.17 \pm 7.31 (N=12)	$\Delta=5.00$, p=0.0003***	$\Delta=4.17$, p=0.0052**	$\Delta=9.17$, p=0.0003***
qwen-max	62.08 \pm 12.33 (N=12)	72.08 \pm 8.53 (N=12)	79.58 \pm 9.23 (N=12)	$\Delta=10.00$, p=0.0012**	$\Delta=7.50$, p=0.0000***	$\Delta=17.50$, p=0.0000***
qwq-32b:free	70.83 \pm 10.17 (N=12)	77.67 \pm 9.30 (N=12)	88.42 \pm 6.37 (N=12)	$\Delta=6.83$, p=0.0137*	$\Delta=10.75$, p=0.0000***	$\Delta=17.58$, p=0.0000***
OVERALL	64.08 \pm 15.25 (N=120)	69.07 \pm 16.63 (N=120)	75.20 \pm 15.39 (N=120)	$\Delta=4.99$, p<0.001***	$\Delta=6.13$, p<0.001***	$\Delta=11.12$, p<0.001***

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: **[TODO]**

Justification: **[TODO]**

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: **[TODO]**

Justification: **[TODO]**

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: **[TODO]**

Justification: **[TODO]**

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: **[TODO]**

Justification: **[TODO]**

Table 10: Overall Mean (\pm SD, N) Confidence and Paired Test Results for Confidence Escalation Averaged Across All Experiment Types.

Model	Opening Bet	Rebuttal Bet	Closing Bet	Open→Rebuttal	Rebuttal→Closing	Open→Closing
anthropic/claude-3.5-haiku	67.71 \pm 10.31 (N=48)	72.60 \pm 10.85 (N=48)	77.19 \pm 11.90 (N=48)	$\Delta=4.90$, p=0.0011**	$\Delta=4.58$, p=0.0003***	$\Delta=9.48$, p=0.0000***
anthropic/claude-3.7-sonnet	57.67 \pm 8.32 (N=49)	63.47 \pm 8.16 (N=49)	68.67 \pm 11.30 (N=49)	$\Delta=5.80$, p=0.0000***	$\Delta=5.20$, p=0.0000***	$\Delta=11.00$, p=0.0000***
deepseek/deepseek-chat	58.65 \pm 11.44 (N=48)	63.23 \pm 11.39 (N=48)	64.58 \pm 11.76 (N=48)	$\Delta=4.58$, p=0.0000***	$\Delta=1.35$, p=0.0425*	$\Delta=5.94$, p=0.0000***
deepseek/deepseek-r1-distill-qwen-14b:free	70.09 \pm 14.63 (N=47)	71.06 \pm 15.81 (N=47)	74.17 \pm 15.35 (N=47)	$\Delta=0.98$, p=0.2615	$\Delta=3.11$, p=0.0318*	$\Delta=4.09$, p=0.0068**
google/gemini-2.0-flash-001	44.88 \pm 25.35 (N=48)	51.54 \pm 20.67 (N=48)	53.73 \pm 17.26 (N=48)	$\Delta=6.67$, p=0.0141*	$\Delta=2.19$, p=0.2002	$\Delta=8.85$, p=0.0041**
gemma-3-27b-it	63.33 \pm 10.42 (N=48)	70.52 \pm 15.52 (N=48)	79.79 \pm 13.07 (N=48)	$\Delta=7.19$, p=0.0008***	$\Delta=9.27$, p=0.0000***	$\Delta=16.46$, p=0.0000***
gpt-4o-mini	68.02 \pm 10.29 (N=48)	72.75 \pm 13.65 (N=48)	78.33 \pm 9.59 (N=48)	$\Delta=4.73$, p=0.0131*	$\Delta=5.58$, p=0.0006***	$\Delta=10.31$, p=0.0000***
o3-mini	67.40 \pm 12.75 (N=48)	71.56 \pm 13.20 (N=48)	73.62 \pm 14.70 (N=48)	$\Delta=4.17$, p=0.0000***	$\Delta=2.06$, p=0.0009***	$\Delta=6.23$, p=0.0000***
qwen-max	60.83 \pm 17.78 (N=48)	69.50 \pm 13.48 (N=48)	75.77 \pm 12.53 (N=48)	$\Delta=8.67$, p=0.0000***	$\Delta=6.27$, p=0.0000***	$\Delta=14.94$, p=0.0000***
qwq-32b:free	67.92 \pm 12.62 (N=48)	73.75 \pm 15.23 (N=48)	78.48 \pm 17.44 (N=48)	$\Delta=5.83$, p=0.0000***	$\Delta=4.73$, p=0.0000***	$\Delta=10.56$, p=0.0000***
GRAND OVERALL	62.62 \pm 15.91 (N=480)	67.98 \pm 15.57 (N=480)	72.42 \pm 15.71 (N=480)	$\Delta=5.36$, p<0.001***	$\Delta=4.44$, p<0.001***	$\Delta=9.80$, p<0.001***

Table 11: Count of Models with Statistically Significant Confidence Escalation per Transition and Experiment Type (One-sided Paired t-test, $p \leq 0.05$).

Experiment Type	Open→Rebuttal	Rebuttal→Closing	Open→Closing
cross_model	6/10	8/10	9/10
informed_self	4/10	1/10	6/10
public_bets	7/10	8/10	10/10
self_debate	7/10	7/10	8/10

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [TODO]

Justification: [TODO]

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer, etc.) necessary to understand the results?

Answer: [TODO]

Justification: [TODO]

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [TODO]

Justification: [TODO]

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [TODO]

Justification: [TODO]

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [TODO]

Justification: [TODO]

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [TODO]

Justification: [TODO]

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pretrained language models, image generators, or scraped datasets)?

Answer: [TODO]

Justification: [TODO]

1158 **12. Licenses for existing assets**

1159 Question: Are the creators or original owners of assets (e.g., code, data, models), used in

1160 the paper, properly credited and are the license and terms of use explicitly mentioned and

1161 properly respected?

1162 Answer: **[TODO]**

1163 Justification: **[TODO]**

1164 **13. New assets**

1165 Question: Are new assets introduced in the paper well documented and is the documentation

1166 provided alongside the assets?

1167 Answer: **[TODO]**

1168 Justification: **[TODO]**

1169 **14. Crowdsourcing and research with human subjects**

1170 Question: For crowdsourcing experiments and research with human subjects, does the paper

1171 include the full text of instructions given to participants and screenshots, if applicable, as

1172 well as details about compensation (if any)?

1173 Answer: **[TODO]**

1174 Justification: **[TODO]**

1175 **15. Institutional review board (IRB) approvals or equivalent for research with human**

1176 **subjects**

1177 Question: Does the paper describe potential risks incurred by study participants, whether

1178 such risks were disclosed to the subjects, and whether Institutional Review Board (IRB)

1179 approvals (or an equivalent approval/review based on the requirements of your country or

1180 institution) were obtained?

1181 Answer: **[TODO]**

1182 Justification: **[TODO]**

1183 **16. Declaration of LLM usage**

1184 Question: Does the paper describe the usage of LLMs if it is an important, original, or

1185 non-standard component of the core methods in this research? Note that if the LLM is used

1186 only for writing, editing, or formatting purposes and does not impact the core methodology,

1187 scientific rigorousness, or originality of the research, declaration is not required.

1188 Answer: **[TODO]**

1189 Justification: **[TODO]**